

Banner appropriate to article type will appear here in typeset article

Designing wake-supervised latent states for predictive modelling of vortex-gust airfoil interactions

Alberto Solera-Rico^{1,2}, Arnau Miró^{3,4}, Oriol Lehmkuhl⁴ and Carlos Sanmiguel Vila^{1,2}

¹Aerial Platforms Department, Spanish National Institute for Aerospace Technology (INTA), San Martín de la Vega, 28330, Madrid, Spain

²Department of Aerospace Engineering, Universidad Carlos III de Madrid, Leganés, 28911, Madrid, Spain

³TUAREG Research Group, Universitat Politècnica de Catalunya (UPC), Carrer Colom 11, 08222 Terrassa, Spain

⁴Barcelona Supercomputing Centre (BSC), CASE Large-Scale Computational Fluid Dynamics, Plaça d'Eusebi Güell, 1-3, 08034 Barcelona, Spain

Corresponding author: Carlos Sanmiguel Vila, csanvil@inta.es

(Received xx; revised xx; accepted xx)

Extreme vortex-gust encounters of a post-stall airfoil are wake-reorganisation events: the load spike is the surface signature of leading-edge-vortex roll-up and shedding, and the release parameters alone do not fix the wake after impact. Using direct numerical simulations of a NACA 0012 airfoil at $\alpha = 14^\circ$ and $Re = 5000$, perturbed by Taylor vortices spanning gust strength, core diameter and wall-normal offset, we ask which reduced state preserves that wake. At fixed dimension, under one shared probe and predictor, we compare a wake-supervised predictive state, a lift-augmented reconstructive autoencoder (Fukami-Taira lineage), and a proper orthogonal decomposition basis; matched controls separate objective from supervision. Probed from held-out fields sixteen frames after impact, the wake-supervised state reads the wake enstrophy ($R^2 = 0.79$) where the no-wake baselines are at or below the mean predictor; but an objective-free supervised encoder reads it equally well, and a matched reconstructive control reads it comparably under case-disjoint retraining, so the wake supervision, not the objective, supplies the readability. The predictive objective is distinguished dynamically: it gives the most accurate wake rollout, reported as secondary. The mechanism is geometric: the anti-collapse regulariser keeps the predictive and supervised rollouts on the encoded manifold while the reconstructive rollout departs along low-variance directions, and the predictive no-gust shedding orbit stays a single loop where the reconstructive one fragments. Sparse wall pressure recovers the predictive state best from a strictly pre-impact window, a state-recovery property rather than direct wake recovery.

Key words: vortex shedding, low-dimensional models, machine learning

1. Introduction

A discrete vortex gust striking a post-stall airfoil does not merely perturb the force coefficients: it reorganises the separated wake. The lift excursion it produces is the surface signature of leading-edge-vortex roll-up, growth, and shedding, the dominant unsteady mechanism in these encounters (Eldredge & Jones 2019), together with the trailing-edge and wake structures it leaves behind, and it is set by the timing of the vortex impact relative to the airfoil's own shedding, the wall-normal offset of the vortex core, and the gust strength and core size. Such extreme gusts arise wherever the gust velocity is comparable to or larger than the flight speed, for small uncrewed and micro air vehicles in urban canyons, mountainous terrain, and the wakes of buildings and ships (Jones *et al.* 2022; Mohamed *et al.* 2023): the gust ratio G , the ratio of the characteristic gust velocity to the free-stream velocity, then exceeds unity, and conventional gust models do not capture the resulting large, fast load transients (Fukami *et al.* 2024, 2025). Even at the moderate chord Reynolds number $Re = 5000$, the post-stall flow at $\alpha = 14^\circ$ is massively separated, and the gust-induced disturbance interacts nonlinearly with the natural shedding cycle (Fukami *et al.* 2025); the vortical organisation of the encounter is moreover similar across laminar and turbulent regimes (Odaka *et al.* 2026), so this case is representative of the wider extreme-aerodynamic family. Resolving the full parametric envelope by simulation is expensive, so reduced-order models (ROMs) of the encounter-to-encounter and case-to-case variability are an active line of work. A reduced model for such an encounter must do more than reproduce a vorticity snapshot or the instantaneous force: it must retain the wake variables that determine the post-impact transient and organise them into a state on which a forward model can be queried without leaving the data manifold. This is the predictive-state problem that organises the present study.

Reduced-order models of such flows form a ladder of increasing nonlinearity. Energy-optimal linear bases, proper orthogonal decomposition (Berkooz *et al.* 1993) and dynamic mode decomposition (Schmid 2010), are compact and interpretable but limited for strongly nonlinear separation; their Koopman-operator reading (Lusch *et al.* 2018) nonetheless makes a linear latent a principled comparator rather than a strawman. Nonlinear autoencoders compress further at matched latent dimension (Murata *et al.* 2020; Lee & Carlberg 2020), and coupling an autoencoder to a learned latent dynamics forecasts better than dynamics confined to a linear subspace (Linot & Graham 2023). Read as a design ladder, these methods progress from field reconstruction and energy compression, through observable readability, toward forward-state usability, which is the property a reduced state needs to serve a predictor; the extreme-aerodynamic reduced models we build on sit at the nonlinear, observable-augmented rung, and the question we pursue is what carries a state the rest of the way.

Recent ROMs for separated and extreme-aerodynamic flows, increasingly built with machine learning (Brunton *et al.* 2020), share a common ingredient: a reconstruction-trained autoencoder whose bottleneck is shaped by an auxiliary aerodynamic observable. Fukami & Taira (2023) trained a convolutional autoencoder jointly with a lift-prediction head and found that the auxiliary loss reduced the intrinsic dimension of the discovered manifold for extreme vortex-gust flows from roughly ten to three. The same observable-augmentation idea was extended across multi-source turbulent data (Fukami & Taira 2025), and was used by Solera-Rico *et al.* (2024) with a β -variational autoencoder and a transformer ROM on a vortex-gust airfoil dataset of the same family considered here. Across this lineage the reconstruction term anchors the latent geometry to the instantaneous

flow field, and an observable-prediction term orients it along aerodynamically relevant directions. A recurring theme of the geometric-analysis literature, however, is that the latent coordinates of an autoencoder are unique only up to a diffeomorphism, so reconstruction leaves the geometry of the latent largely unconstrained, which complicates any downstream task that relies on that geometry, including dynamics, interpolation, and comparison (Tran *et al.* 2026; Kelshaw & Magri 2024). Observable augmentation solves only one part of the predictive-state problem: it tells the encoder which aerodynamic quantity to represent, but it does not by itself fix whether the latent coordinates are suitable for time advancement. A reduced state that is linearly readable at a single instant may still be unusable once an autoregressive predictor queries the points its own rollout produces.

A parallel and fast-growing line reads these flows through the information their reduced states carry rather than their energy. Transfer-entropy causality among proper-orthogonal-decomposition modes ranks which modes drive the formation of coherent structures (Martínez-Sánchez *et al.* 2023), and an informative/non-informative decomposition splits a field into the part that carries information about a target’s future and a residual that carries none (Arranz & Lozano-Durán 2024). This informative-decomposition idea has very recently been brought to extreme vortex-gust airfoil interactions of the present family, extracting the field structures informative about the future lift (Koshikawa *et al.* 2026; Fukami & Araki 2026), alongside time-dependent modal bases for the same interaction (Zhong *et al.* 2025; Zamani Ashtiani & Fukami 2025) and latent-causality attributions in physical space (Cremades *et al.* 2026). Our study is complementary: rather than mapping the informative or causal modes of a single representation, we ask which encoder objective produces a latent whose forward evolution closes the full set of observables, the wake included, and we use the information lens only as a held-out diagnostic of the resulting latent (§5.1).

A complementary route to a reduced state is to drop the reconstruction term altogether and to train the encoder and a predictor end to end on latent-space forecasting, with an explicit anti-collapse regulariser in place of the decoder. This is the joint-embedding predictive architecture (JEPA) (LeCun 2022): an encoder maps the input to a representation, a predictor advances that representation in time, and the loss never reconstructs the field. The encoder is thereby trained in the same coordinates in which the predictor will later operate, which makes JEPA a useful instrument for testing whether predictive training produces a latent state that remains usable when rolled forward, rather than merely readable when encoded from the DNS field; the predictor also cannot exploit pixel-level shortcuts, because no field-space signal reaches it. The image and video instantiations (Assran *et al.* 2023, 2025) and the recent from-pixels variants that remove the moving-target encoder in favour of a characteristic-function regulariser (Maes *et al.* 2026; Balestrieri & LeCun 2025) or a variance-covariance objective (Sobal *et al.* 2025; Bardes *et al.* 2022) have so far been evaluated on gridworld and toy visual domains. Their behaviour on a low-intrinsic-dimension physics dataset, where the data sit on a roughly five- to ten-dimensional manifold, is uncharacterised.

Reconstruction-trained nonlinear encoders compress aggressively but yield latent dynamics that are harder to forecast than an energy-optimal linear basis, a compactness-versus-forecast tradeoff documented for controlled wake flows in a companion study under review (Solera-Rico *et al.* 2025). We show this tradeoff is a consequence of the reconstruction objective: a predictive latent recovers nonlinear compactness without surrendering forecast stability. The division of labour with that companion study (Solera-Rico *et al.* 2025) is explicit: there, controlled canonical wakes and the compactness-versus-forecast tradeoff; here, the parametric vortex gust, the latent-drift mechanism we identify as the common cure, and its geometric and topological characterisation.

The model we adopt takes no gust parameters: (G, D, Y) enter neither the encoder nor the predictor. The encoder is a pure state map and the predictor advances the latent from its own history alone. The world-model framework of LeCun (2022); Ha & Schmidhuber (2018) motivates this split as an analogy and not a property we claim, since the latent dynamics function is the object a controller would plan against, and a recent gust-airfoil instance couples a physics-augmented autoencoder to a latent dynamics model (Liu *et al.* 2025). We treat that reading as motivation: the rollout is shown as a consistent latent trajectory rather than advanced as a validated forecast, and its interventional or causal interpretation is left for future work (§5). What this leaves is the property our study turns on: an established reconstruction-trained autoencoder and an energy-optimal proper orthogonal decomposition basis, lacking wake supervision, do not keep the dynamically relevant wake observables linearly readable, and even a wake-readable state is not automatically usable once a forward model advances it. The drift mechanism of §4.3 and the division of labour we identify are the empirical consequences.

The question this paper addresses is therefore not which encoder reconstructs the flow best, nor even which produces a reduced state from which the wake is readable at a single instant, but which produces a state that is also usable once a forward model advances it. We treat this as a fluid-mechanics question with a geometric answer, and find a division of labour: wake-observable supervision makes the wake readable, an anti-collapse regulariser keeps the latent in distribution under rollout, and the predictive objective makes the supervised state dynamically trackable. The joint-embedding predictive architecture is the instrument we use to test this principle, not the subject of the paper. We make three contributions.

First, we formulate the vortex-gust encounter as a predictive-state benchmark. The primary endpoint is the wake enstrophy at a fixed post-impact readout ($H = 16$), because it summarises the leading-edge-vortex roll-up and wake reorganisation that the scalar force can miss; all encoders, a predictive one (JEPA), a reconstructive one in the Fukami lineage, and a POD basis, are compared at matched latent dimension under a single shared probe and autoregressive predictor (§3). The question is then not field-reconstruction quality but whether the reduced state carries the wake variables a forward model needs (§4.1).

Second, we separate the ingredients of a usable predictive state (§4.2). Wake-observable supervision supplies the instantaneous readability: on held-out fields the wake-supervised states make the wake enstrophy linearly readable ($R^2 = 0.79$) where the no-wake reconstructive and linear baselines are at or below zero, but an objective-free supervised encoder reads it as well, and a matched reconstructive control reads it comparably once retrained out of sample, so the readability is the supervision's, not the objective's. Anti-collapse regularisation supplies an in-distribution geometry: the supervision-only and predictive rollouts both stay on the encoded manifold, where the reconstructive rollout departs along directions of negligible encoded variance (§4.3). The predictive objective supplies the dynamic usability: its rollout stays closer to the encoded simulation trajectory and forecasts the wake better than the objective-free and reconstructive alternatives.

Third, we interpret the resulting state physically and observationally (§4.4, §4.7). The predictive state keeps the undisturbed shedding orbit as a single loop under a metric that removes coordinate scale, carries the wake response across the tested readout horizons, and is the most recoverable of the three from sparse pre-impact wall pressure. These results motivate the state as an input to a future estimator or controller; the present paper does not claim a validated forecast or a closed-loop controller.

2. Flow configuration and data

2.1. Vortex gust-airfoil interaction

We study a NACA 0012 airfoil at angle of attack $\alpha = 14^\circ$ and chord Reynolds number $Re = 5000$, perturbed by a discrete vortex gust modelled as a spanwise-oriented Taylor vortex released upstream of the leading edge. The data analysed here are our own direct numerical simulations of this configuration, computed with the GPU-enabled spectral-element solver SOD2D (Gasparino *et al.* 2024) with no subgrid-scale model, so that all scales of the separated wake are resolved at this Reynolds number. The solver is run in the low-Mach regime in which the flow is effectively incompressible. The configuration and its physical characterisation follow Fukami, Smith & Taira (2025), which is the reference for the physics, not the source of the data. The vortex has the azimuthal velocity profile

$$u_\theta(r) = u_{\theta,\max} \frac{r}{R} \exp\left[\frac{1}{2}\left(1 - \frac{r^2}{R^2}\right)\right], \quad (2.1)$$

which peaks at the core radius $r = R$ (Taylor 1918), and is parametrised by three scalars: the signed gust ratio $G \equiv u_{\theta,\max}/u_\infty$, the core diameter $D \equiv 2R/c$, and the wall-normal offset Y of the vortex centre, the displacement normal to the chord (the laboratory coordinates are rotated by the angle of attack to define Y). The released-data case identifiers encode the solver's opposite sign convention, $s = -G$, so a case labelled G+4.00 in the archive corresponds to $G = -4$ here. The profile and the (G, D, Y) sampling follow the configuration of Fukami *et al.* (2025). The training envelope spans $|G| \in \{0.5, 1.0, 1.5, 2.0, 3.0\}$ with $G = 0$ as the undisturbed reference, $D \in \{0.5, 1.0, 1.5\}$ chords, and $Y \in [-0.4, +0.4]$; the value $|G| = 4$ is held out as an extrapolation boundary, for reasons that are physical and that we return to below. Convective time t is measured in chords and set to zero when the vortex centre reaches the leading edge.

Without a gust, the $\alpha = 14^\circ$ wake at this Reynolds number is massively separated and sheds periodically; in dynamical terms the undisturbed flow is a limit cycle. A gust encounter is then a transient excursion from this limit cycle and back, and it is helpful to read it in stages (figure 1) (Smith *et al.* 2024; Fukami *et al.* 2025; Odaka *et al.* 2026). The vortex first induces a strong wall-normal vorticity flux at the leading edge, the airfoil experiences a sharp lift excursion as a leading-edge vortex (LEV) forms, the LEV rolls up and sheds together with a trailing-edge structure as the load reaches its extremum, and the wake then relaxes back toward the baseline shedding. The sign of G sets whether the first lift excursion is positive or negative, Y sets the side and strength of the impingement, D sets how much of the encounter is dominated by the vortex core, and the impact timing sets the phase of the baseline cycle at which the disturbance arrives. The six observables defined in §2.2 are quantitative summaries of exactly this sequence.

A consequence of the source characterisation (Fukami *et al.* 2025) is central to how we read the extrapolation case. For $|G| \leq 3$ the encounter is primarily two-dimensional, so a mid-plane slice of the spanwise vorticity carries the relevant large-scale dynamics; for $|G| \geq 4$ the interaction becomes genuinely three-dimensional, with fine-scale instabilities developing during impact. We confirm this from the full three-dimensional vorticity fields: the fraction of the spanwise-vorticity (ω_z) enstrophy not captured by its spanwise mean, the content a mid-plane slice discards, is roughly constant across the training range ($|G| \leq 3$, post-impact wake median ≈ 0.20) and rises about threefold at $|G| = 4$ (≈ 0.56 , over the four available cases). Because the encoder in this work consumes a single mid-plane vorticity field (§2.2), the held-out $|G| = 4$ case is not merely one parameter step beyond the training range; it is a regime in which the quantity the encoder observes no longer determines the




Figure 1. The vortex gust as a wake-reorganisation event, shown as a gust-strength sweep at fixed core diameter $D = 0.5$ and wall-normal offset $Y = -0.1$ (the first encounter of each case). (a) The lift coefficient C_L against convective time t/c for the four cases, with $t/c = 0$ the gust-release instant and impact near $t/c = 2$; the undisturbed baseline ($G = 0$) is grey, and the load excursion deepens and widens monotonically with the gust strength $|G|$. (b) The mid-plane spanwise vorticity ω_z for each case (rows, $G = 0, -0.5, -1.5, -4$) at four convective times (columns, $t/c = 1, 2, 3, 4$), in raw physical units with the colorbar clipped at ± 3 (the intense cores saturate; the wake structure away from them lies within the range), and the airfoil drawn as a filled polygon. The strong gust injects a vortex that reorganises the wake through impact, while the weak cases perturb the baseline shedding only slightly; $G = -4$ is the $|G| = 4$ extrapolation case, shown here only as a physical illustration. The load spike and the wake reorganisation are the surface and field signatures of the same event, and the release parameters alone do not fix where the wake sits afterwards, which is why the wake, not only C_L , is the endpoint a reduced state must carry.

191 true dynamics. We therefore treat $|G| = 4$ as a physical observability boundary rather than
 192 as a claim of parametric generalisation.

2.2. Data, fields, and observables

194 Each simulation case is one long run partitioned into encounters of 120 cached frames at
 195 $\Delta t = 0.05 t/c$, one encounter per gust release. The simulations use a spectral-element
 196 discretisation of polynomial order $p = 4$, and the fields are cached on a structured
 197 grid of $192 \times 96 \times 32$ points spanning $x/c \in [-1.5, 4.5]$ and $y/c \in [-1.5, 1.5]$ with
 198 a spanwise extent of approximately one chord; this analysis subdomain matches the data-
 199 driven subdomain of the source characterisation (Fukami *et al.* 2025). The lift and drag
 200 coefficients are $C_L = F_L / (\frac{1}{2} \rho u_\infty^2 c)$ and $C_D = F_D / (\frac{1}{2} \rho u_\infty^2 c)$, obtained by integrating the
 201 pressure and viscous stresses over the airfoil surface. The remaining DNS solver-resolution
 202 details required for full reproducibility, the free-stream Mach number, the domain and
 203 spanwise extent, the element and solution-point counts, the near-wall spacing, the time
 204 step and CFL, the gust-release station, and the grid and time-step sensitivity, are listed as
 205 an author-fill checklist in table 1. The dataset comprises 85 cases partitioned into training,
 206 validation (the held-out encounters of the training cases, denoted test_a), an in-distribution
 207 test set (test_b), and an out-of-distribution test set (test_c). The in-distribution test set
 208 is stratified to span the five gust strengths, the three core diameters, and the two source
 209 groups (figure 2); test_c is the $|G| = 4$ extrapolation set. These 85 cases are separated
 210 in preprocessing into 382 encounter windows of 120 frames each, one per gust release:
 211 229 training, 87 validation (test_a, the held-out last encounters of the training cases), 42
 212 test_b, and 24 test_c. The undisturbed reference case is retained in training. Behind these
 213 encounter counts are far fewer distinct cases: the 42 test_b encounters come from only 10
 214 held-out cases (five interior and five boundary tier) and the 24 test_c encounters from 4
 215 cases, with most cases contributing four to six encounters that share a baseline shedding
 216 history. Encounters of the same case are therefore not independent, so encounter-level
 217 intervals overstate the effective sample size; the headline wake comparison is reported
 218 additionally with a case-clustered bootstrap and a case-level paired test (Appendix A). The
 219 partition is fixed once and used identically by every model and baseline.

220 For each frame the cache stores the mid-plane spanwise vorticity ω_z (192, 96), the span-
 221 averaged wall pressure p_{wall} (192), and the lift and drag coefficients C_L , C_D . The vorticity is
 222 normalised by a fixed pipeline: a spatial mask removes the solid and one adjacent cell layer
 223 to suppress the leading-edge finite-difference artefact, a per-encounter high-percentile clip
 224 caps spikes, and the field is scaled by the training standard deviation with no mean shift,
 225 preserving the vorticity antisymmetry. All representation and probe losses are computed
 226 in this normalised space; unnormalised values are used only for physical-units metrics and
 227 figures. Appendix A repeats the primary wake probe under no amplitude clip and under
 228 a single training-pooled leak-free threshold, and the family ordering is unchanged, so the
 229 per-encounter clip does not drive the result.

230 We evaluate forward closure against six physical observables that together span body-
 231 force, integrated-flow, and spatially distributed wake diagnostics. The lift and drag
 232 coefficients C_L and C_D are the body forces. The wall-normal impulse I_y is an integrated-
 233 flow quantity coupled to the wake-vortex dynamics. The wake enstrophy and the positive
 234 and negative circulation are spatially distributed wake observables: they integrate (signed)
 235 vorticity over the wake region and are the most demanding to forecast, because they are
 236 sensitive to the redistribution of vorticity that the LEV roll-up and shedding produce. The
 237 choice is deliberate: the wake observables can change sharply while the scalar forces change
 238 little, so they are the part of the encounter most likely to separate a representation that has
 239 captured the gust-vortex physics from one that has only captured its force signature. On

Quantity	Value	Why a referee needs it
Free-stream Mach number (or incompressible-limit confirmation)	[PENDING:]	Confirms the low-Mach formulation is effectively incompressible at $Re = 5000$.
Computational domain and spanwise extent L_z/c	[PENDING:]	Bounds domain blockage and confirms the span resolves the three-dimensional wake.
Element and solution-point counts	[PENDING:]	Fixes the spatial resolution of the spectral-element discretisation.
Minimum wall-normal spacing $\Delta n_{\min}/c$ and wall units Δn^+	[PENDING:]	Shows the near-wall layer is resolved without a wall model.
Time step $\Delta t u_\infty/c$ and maximum CFL	[PENDING:]	Establishes the temporal resolution and stability of the integration.
Gust-release station x_0/c	[PENDING:]	Sets the upstream distance over which the Taylor vortex develops before impact.
Grid and time-step sensitivity check (baseline and a representative gust case)	[PENDING:]	Demonstrates that the reported statistics are converged.

Table 1. DNS solver-resolution details required for full reproducibility (§2.2). The data are direct numerical simulations with no subgrid-scale model, so all scales are resolved and no large-eddy closure is involved. Values marked pending are owned by the simulation collaborators and are filled from `paper/dns_metadata.yaml` before submission.

these grounds we fix the wake enstrophy in advance as the single primary endpoint of the closure comparison; the other five observables are reported as secondary and descriptive. This pre-registration is what the family-wide multiplicity correction in Appendix A is applied against.

The six observables are defined on each frame of the masked vorticity field ω_z described above. The lift and drag coefficients C_L, C_D are as defined earlier in this section. On the wake window $\Omega_w = \{(x, y) : x/c \in [0.5, 4], |y/c| \leq 1\}$ the wake enstrophy is

$$E_w(t) = \int_{\Omega_w} \omega_z^2 dA, \quad (2.2)$$

and the signed wake circulations integrate the vorticity over its strongly rotational parts,

$$\Gamma^+(t) = \int_{\Omega_w, \omega_z > \omega_c} \omega_z dA, \quad \Gamma^-(t) = \int_{\Omega_w, \omega_z < -\omega_c} \omega_z dA, \quad (2.3)$$

with threshold $\omega_c = 1$ in the cleaned vorticity units. The signed circulations at $\omega_c \in \{0.5, 1, 2\}$ are nearly collinear (pairwise Pearson well above 0.9 across held-out frames) and the predictive latent’s circulation closure advantage over the reconstructive baseline is essentially unchanged across them, so the closure does not depend on this threshold. The wall-normal hydrodynamic impulse per unit span is

$$I_y(t) = \int_{\Omega} x \omega_z dA \quad (2.4)$$

over the full cached field Ω . Every case uses the same fixed window and mask, and values are reported in the cache vorticity units integrated over chord-normalised area. The forecast horizon $H = 16$ is the sixteenth cached frame after impact, $t = 16 \Delta t = 0.8 c/u_\infty$.

Two definitional points bound how these observables are read. The impulse I_y is the mid-plane two-dimensional vorticity impulse, a diagnostic of wake-vortex transport, and not the impulse-theorem lift: the mid-plane slice omits the bound circulation, so on the




Figure 2. Parameter-space sampling of the 85 cases in the (G, D, Y) gust envelope, shown as one isometric view and coloured by split: training (229 encounters), the stratified in-distribution test set (test_b, 42 encounters, split into interior and boundary tiers), and the out-of-distribution extrapolation set (test_c, 24 encounters). The translucent plane marks the $|G| = 4$ extrapolation boundary and the grey points on the floor are the projection of every case, which shows that training mass concentrates near $Y/c = 0$, the reason the wall-normal offset is the least well resolved parameter (§4.6).

simulation dI_y/dt does not track the lift (Pearson correlation near zero on test_b, far from the value near -1 that a planar impulse-lift relation would give), and no closure result below should be read as recovering the impulse-theorem lift. More broadly, the wake and flow observables (the enstrophy, the two circulations, and I_y) are computed on the same mid-plane slice the encoder consumes, so they omit the out-of-plane content quantified in §2.1 (roughly a fifth of the ω_z enstrophy in distribution), whereas C_L and C_D are true three-dimensional surface-integrated forces. The closure therefore compares mid-plane wake diagnostics against three-dimensional forces, and the mid-plane proxy is weakest at the $|G| = 4$ extrapolation boundary, where the out-of-plane fraction rises substantially.

Table 2. Encoder and predictor configuration. The model is unconditional: no gust parameters $\mathbf{c} = (G, D, Y)$ enter the encoder or the predictor. Parameter counts are for the $d = 64$ configuration; two small auxiliary heads read observables off the encoder output during training (a lift head, 4×10^3 parameters, and an 80-dimensional wake-spectrum head, 3.5×10^4 parameters; §3.1).

	Encoder E_θ	Predictor P_ϕ
role	field $\rightarrow$ latent (unconditional)	latent dynamics (unconditional)
input	$\omega_z \in \mathbb{R}^{192 \times 96}$	latent sequence $\mathbf{z}_{1:t}$
backbone	3 conv. stages + 6-layer ViT	6-layer autoregressive transformer
width / heads	256 / 8	384 / 16
tokens / context	$24 \times 12 = 288$ spatial	≤ 32 temporal (causal)
conditioning	none	none (RoPE temporal positions, causal mask)
latent read-out	class token $\rightarrow$ linear $\rightarrow$ BatchNorm	next-step latent
parameters	6.7M	16.2M

3. Predictive latent representation and evaluation protocol

3.1. Model and training objective

The model takes no gust parameters: neither the encoder nor the predictor receives (G, D, Y) . The predictive encoder E_θ maps the mid-plane vorticity field $\omega_z(192, 96)$ to a latent $\mathbf{z} \in \mathbb{R}^d$ at $d \in \{16, 32, 64\}$. It is a hybrid network: three convolutional downsampling stages produce a $24 \times 12 \times 256$ feature map, a six-layer vision transformer processes the resulting tokens, and the class-token output is projected to $\mathbf{z}$ through a single linear layer with a batch-normalisation boundary, which the characteristic-function regulariser below requires. The encoder is therefore a pure state map, and this is the property that lets the latent itself serve as the observable a controller would have access to, since at deployment the gust parameters are not known a priori (§5).

The predictor P_ϕ advances the latent in time. It is a six-layer autoregressive transformer with rotary positional embeddings and a causal mask, and it advances the latent from its own history alone, so the model is unconditioned end to end. Its internal structure is given in Appendix A. Figure 3(a) shows the encoder and predictor, and panel (b) sets out the matched-predictor evaluation protocol used throughout; the contrast between the predictive training signal and the reconstructive one, in which the reconstructive autoencoder passes the field at time t through an encoder and a decoder and incurs a field-space error whereas the predictive model encodes the fields at t and $t + \Delta t$ with shared weights and incurs a latent-space error between the predicted and observed target embeddings, with no field-space loss, is drawn in Appendix A (figure 13). Table 2 summarises the two networks.

The predictive encoder and predictor are trained jointly on a one-step latent-prediction term, a scheduled-sampling rollout term that exposes the predictor to its own errors, and an anti-collapse regulariser that prevents the encoder from collapsing to a constant. We use the characteristic-function regulariser of Balestrieri & LeCun (2025); Maes *et al.* (2026), which matches the latent distribution to an isotropic Gaussian through a small number of random projections, with an automatic fall-back to a variance-covariance objective (Bardes *et al.* 2022) should the participation ratio of the latent indicate collapse. Its configuration is given in Appendix A.

Two auxiliary heads read aerodynamic observables off the encoder output during training: a lift head that maps $\mathbf{z}_t$ to $C_L(t)$, and a wake head that maps $\mathbf{z}_t$ to an 80-dimensional signed wake spectrum, both trained with small weights jointly with the representation and neither part of the latent-prediction loss. These heads are adopted from the observable-augmentation lineage (Fukami & Taira 2023, 2025). They are also,




Figure 3. The predictive model and how it is evaluated. (a) The encoder maps the mid-plane vorticity field to a latent $\mathbf{z}_t$, and the autoregressive predictor advances the latent from its own history; no gust parameters enter the encoder or the predictor. The two auxiliary observable heads (lift and wake) read $\mathbf{z}_t$ during JEPA training with small weights and are not part of the latent-prediction loss; they are omitted from the schematic for clarity. The visualisation decoder is trained only after the encoder is frozen and is used solely for the diagnostic decodes of §4.6, never in the predictive loss. (b) Evaluation applies the same predictor and probe family g to every encoder family: a probe on the simulation-encoded latent measures representational closure, the headline, and the same probe on the rolled-out latent $\hat{\mathbf{z}}_{t+H}$ measures forecast closure, reported as secondary.

however, the reason the present predictive encoder is not a clean isolation of the training objective: the reconstructive baseline below carries a lift head but not a wake head, and the two encoders differ in architecture as well. The headline comparison is therefore between encoder *families* as configured, and the controls of §4.2 separate the objective from the architecture and the wake supervision.

3.2. Matched-predictor closure protocol

The comparison rests on a single principle: every baseline is the same pipeline, and only the encoder differs. Each encoder produces a latent trajectory $\mathbf{z}_{1:T}$ from the vorticity fields, and an identical probe family maps a latent to a physical observable for evaluation (figure 3(b)). The probe family is held fixed across the three encoder families, and the predictor is shared by all of them. The latent dimension is matched within each family: JEPA at $d \in \{16, 32, 64\}$, the reconstructive autoencoder at $d \in \{3, 32, 64\}$ to include both the published $d = 3$ recipe and matched-capacity variants, and POD at $d \in \{16, 32, 64\}$. The predictive family is additionally swept to $d \in \{4, 8\}$ and crossed with a conditioned predictor variant in §4.2. This isolation is the only fair way to attribute representation quality to the encoder rather than to incidental differences in the readout. The reconstructive autoencoder uses its best-documented configuration, with ReLU activations, group normalisation, and a future-lift head at horizons $\{8, 16, 24\}$, rather than the strict published variant (tanh, no normalisation, a current-lift head), which we found gives a worse parametric probe; the comparison is therefore against a tuned baseline, not a hobbled one.

The headline figure of merit is representational closure: a fixed linear probe is applied to the latent encoded directly from a held-out field, and we ask whether it recovers each of the six observables, the wake in particular. This isolates what the encoder carries, with no rollout. We report it as the held-out coefficient of determination R^2 on test_b and test_c, with the per-observable mean absolute error alongside, and bootstrap 95% confidence intervals over encounters. As a secondary, confirmatory measure we also report forecast closure: the predictor is rolled recursively from the pre-impact context to $H = 16$, and the same probe maps the predicted latent to each observable. A small forecast error means the encoder produces a latent whose dynamics, propagated by the predictor, retain the information the probe needs; a large error means either that the encoder discarded that information or that the predictor lost it during rollout, a distinction the drift diagnostic below resolves. The representational closure is in table 4(a) and the forecast closure in table 4(b). We do not headline a forecast- R^2 -versus-horizon curve; the rollout is reported to show that the latent stays usable, not as a validated forecast.

3.3. Diagnostics, controls, and uncertainty

Two diagnostics turn the closure result into a mechanism. The first is the latent drift, which addresses the fact that a predictor with low one-step error need not produce a useful long rollout, because each prediction is fed back as the next input and small geometric errors accumulate. We quantify how far a rollout has left the training manifold by the Mahalanobis distance of the rolled-out latent at horizon H , measured against the distribution of simulation-encoded latents for the same encoder, normalised by the Mahalanobis distance of the simulation-encoded latent itself. A ratio near one means the rollout stays within the natural spread of the training latents; a large ratio means it has gone out of distribution, so that the probe is then queried outside its support.

The second diagnostic is a parameter-only floor, which carries particular force here because the model never sees the gust parameters: it bounds what the parameters alone could explain, so that any closure above it is genuinely supplied by the latent state. We compute it with a kernel-ridge regressor that maps $\mathbf{c} = (G, D, Y)$ directly to each observable at $H = 16$, with no latent input, fit on training and evaluated on the held-out splits. On wake enstrophy this floor is low, so the latent's wake closure is not something the parameters could have supplied. At the impact frame the same parameter floor is high, because the just-released gust nearly determines the instantaneous state; the gap opens as the wake evolves (§4.1).

Table 3. Encoder families compared under the matched-predictor protocol of §3.2. The model takes no gust parameters: none enter the encoder or the predictor. All three families share one autoregressive transformer predictor, training recipe, and probe family; only the encoder differs. “Aux. heads” lists the observable supervision attached to the encoder. The architecture and auxiliary-head columns differ across families, which is why the headline result is attributed to the predictive *family* and isolated only by the controls of §4.2.

Family	Objective	Encoder	Latent dim. d	Aux. heads
JEPA (this work)	predictive (latent)	CNN + ViT	16, 32, 64	lift + wake
Reconstructive AE	reconstructive (field)	CNN	3, 16, 32, 64	lift
POD	linear projection	none	16, 32, 64	none

Table 4. Held-out closure at $H = 16$ on the $n = 42$ test_b encounters, at matched $d = 64$, in two modes. (a) Representational closure, the headline: the coefficient of determination R^2 when a fixed linear probe is applied to the simulation-encoded latent (the latent encoded directly from the held-out field), larger is better. (b) Forecast closure, reported as confirmatory: the R^2 of the full-context rollout against the simulation observable (the variance baseline is the held-out split mean, so $R^2 < 0$ is worse than predicting that mean). The representation is where the families separate: the predictive latent renders the wake enstrophy linearly readable ($R^2 = 0.79$) where the reconstructive (Fukami) and linear (POD) latents are at or below zero. The forecast panel points the same way and is read as secondary; the per-encounter paired test is in table 8 and full probe values with bootstrap CIs in Appendix A. The β -VAE representation row uses the matched-recipe encoder; its single forecast row is rolled out from the faithful-recipe predictor, the only β -VAE rollout available, and is read accordingly.

(a) Representational closure: held-out R^2

	C_L	C_D	I_y	wake-enst.	circ. pos.	circ. neg.	mean
predictive (JEPA)	0.86	0.69	-0.21	0.79	0.72	0.78	0.61
reconstructive (Fukami)	0.56	0.23	0.13	-0.25	0.25	0.03	
β -VAE	0.69	0.52	0.05	0.53	0.54	0.44	0.46
linear (POD)	0.68	0.64	-0.69	-0.16	-0.15	0.32	

(b) Forecast closure (secondary): full-context rollout R^2

	C_L	C_D	I_y	wake-enst.	circ. pos.	circ. neg.
predictive (JEPA)	0.45	0.41	0.25	0.43	0.39	0.74
reconstructive (Fukami)	0.56	0.12	0.08	0.22	0.57	0.22
β -VAE	0.66	0.28	0.30	-0.07	0.29	0.12
linear (POD)	0.83	0.72	-0.02	-0.60	-0.18	0.45

Because the predictive encoder differs from the reconstructive baseline in objective, architecture, and auxiliary supervision at once, the comparison is accompanied by controls that vary one factor at a time: a crossing of objective with architecture at matched auxiliary heads, an ablation of each auxiliary head, and a conditioned predictor variant that adds $\mathbf{c}$ (§4.2). Every held-out R^2 and mean absolute error reported below carries three independent uncertainty signals, following the reporting protocol fixed with the data split: a 2000-resample bootstrap over the held-out encounters, the standard deviation across three encoder retrains from independent seeds, and a five-fold cross-validation over cases on the probe read-out. Their full specification is in Appendix A.

4. Results

4.1. Representational wake closure and the parameter-only floor

The model comparison is meaningful only if the endpoint is physically discriminating, so we begin with the wake itself. At impact the release parameters largely fix the immediate force excursion; by the readout horizon, sixteen frames after impact and past the leading-

edge-vortex shedding, those same parameters no longer fix the downstream wake enstrophy (figure 1), as the parameter-only floor below makes quantitative. This separation, between the force that the gust sets and the wake that the evolving separated flow sets, is why the wake enstrophy is the primary endpoint.

Table 3 summarises the three encoder families and the protocol. We distinguish two questions. Does the latent, encoded directly from a held-out field, render each observable readable by a fixed linear probe (representation quality, the headline, table 4(a))? And does the predictor, rolled out from the impact frame, track it (forecast quality, reported as secondary, table 4(b))? The first isolates the encoder; the second tests the encoder and the learned dynamics together. The result is established at the representation level, where the wake is the discriminator, and the rollout is shown as consistent confirmation rather than the load-bearing test.

The wake is where the families separate. Probed directly from the held-out field sixteen frames after impact, the predictive latent renders the wake enstrophy linearly readable at $R^2 = 0.79$ (figure 4), while the reconstructive autoencoder and the linear basis are at -0.25 and -0.16 , both below the predict-the-mean floor. A nonlinear (kernel-ridge) probe narrows this gap, recovering part of the wake from the matched-head reconstructive control (§4.2), so the claim is that the predictive latent renders the wake *linearly decodable*, not that the reconstructive encoding discards it irrecoverably. The established reconstructive and linear recipes compared here carry no wake-observable head, and the separation is against that: a matched reconstructive control that does carry the head reads the wake comparably once each case is held out of training (§4.2). A β -VAE lineage baseline that also carries the wake head (table 4) reads part of the wake at the representational endpoint (0.53) but does not preserve it under the shared rollout (-0.07), the same pattern. The representational readability is therefore supplied by the wake supervision, and the predictive objective's robust distinction is dynamical, established in §4.2 and §4.3: it reads the wake consistently across retrains, forecasts it from its own rollout, and keeps that rollout on the encoded manifold. The forces and the circulations are read off the same latent cleanly, with C_L at $R^2 = 0.86$, the negative and positive circulations at 0.78 and 0.72, and the drag at 0.69; the mean over the six observables is 0.61. The one weak observable is the wall-normal impulse I_y ($R^2 = -0.21$), a mid-plane vorticity diagnostic that even the energy-optimal linear basis captures only modestly (§2.2), so its weakness is a property of the observable rather than of the predictive latent. The same ordering holds on the training, validation, and extrapolation (test_c) splits reported below.

The rollout points the same way and is read as confirmatory, not as the headline (table 4(b) and figure 16). A matched transformer predictor is trained on each family's latents under one protocol, and rolled out from the full pre-impact context; probing the wake enstrophy at $H = 16$, the forecast R^2 is 0.43 for the predictive latent, ahead of 0.22 for the reconstructive autoencoder and -0.60 for the linear basis. The forecast ordering is the same as the representational one but compressed, which is why we lead with the representation: the families separate most cleanly at the encoder, and the rollout confirms rather than carries the result. The forecast advantage is not statistically separable under case-level clustering (§4.3 and table 8), so we present it as consistent tracking, not as a forecast claim. On the extrapolation set (test_c, $G = -4$) the predictive wake-enstrophy forecast stays the strongest of the three while the integrated forces collapse, consistent with the three-dimensional observability boundary of §2.1.

The representational wake separation is the numerical core of the result. At the independent-case level the effect is positive but not uniform: the paired direct-encoding margin (table 8) is 33.56 in wake-enstrophy units with a case-clustered 95% interval $[7.79, 58.72]$, and the per-case slopegraph (figure 5) gives the predictive latent the lower




Figure 4. Representational closure of the predictive latent at $H = 16$ after gust impact, the headline result. (a) Held-out R^2 of a fixed linear probe applied to the simulation-encoded latent (no rollout) for each of the six observables, on the in-distribution test set (test_b, $n = 42$); the dashed line marks the mean over the six and the one weak observable, the mid-plane impulse I_y , is shaded light. (b) The wake-entropy R^2 across families: the predictive latent (0.79) against the reconstructive autoencoder (-0.25) and the linear basis (-0.16), each at $d = 64$. The wake is the discriminator, and only the predictive latent renders it linearly readable.

Table 5. Parametric (model-free) floor at $H = 16$ after impact. A cross-validated ridge regressor maps the gust parameters $\mathbf{c} = (G, D, Y)$ directly to each observable, with no latent input, fit on the training cases and read out on the in-distribution test split under the closure protocol ($R^2 = 1 - \text{SSE}/\text{SST}$ about the held-out mean). This is the model-free baseline: the predictive latent beating it (last column, the representational closure of table 4(a)) shows the representation carries flow state the parameters alone do not supply. On wake entropy the floor is negative ($R^2 = -0.18$): the gust parameters alone do worse than the held-out mean sixteen frames after impact, because the wake is then governed by the unfolding dynamics rather than by the release condition, whereas the predictive latent reads it at 0.79. On the forces, which track the gust forcing more directly, the floor is high.

Observable	parametric floor (test_b R^2)	predictive repr. (test_b R^2)
C_L	0.61	0.86
C_D	0.54	0.69
I_y	0.25	-0.21
wake-entropy	-0.18	0.79
circ. pos.	-0.05	0.72
circ. neg.	0.23	0.78

error on seven of the ten held-out cases. Across the full twelve-test closure matrix the pre-registered encounter-level Holm family retains the wake-entropy advantage, together with the lift and the negative circulation; the stricter case-level Holm family is underpowered on ten cases and is not used to promote the secondary observables. We therefore state the primary claim as a pre-specified paired wake-entropy result with the case spread shown explicitly, while the forecast rows and the non-primary observables remain descriptive. All six observables are reported in both modes, not a selected subset, and the ordering is not uniform: the predictive latent leads on the spatially distributed wake observables (entropy and negative circulation), which is precisely the paper's claim. This separation is confirmed by the mechanistic evidence of §4.3 (the drift departure spectrum and the no-gust limit-cycle topology), where the predictive and reconstructive families separate consistently.

The parametric floor (table 5) confirms that the wake closure is not something the gust parameters supply on their own. Mapping $\mathbf{c} = (G, D, Y)$ to each observable at $H = 16$




Figure 5. Per-case wake-entropy closure, the honest case-level view of the headline result. Each line is one held-out test_b case (ten in all), connecting the reconstructive autoencoder’s per-case mean absolute error (left) to the predictive latent’s (right) at $H = 16$; a downward line favours the predictive latent. The predictive latent gives the lower error on seven of the ten cases (median per-case improvement positive, case-level signed-rank significant), with three cases favouring the reconstructive baseline: the advantage is real at the case level but not uniform, which is why the case-level multiplicity correction is not passed and the claim is stated at encounter level with this spread shown.

with cross-validated regressors, the floor on wake entropy is negative ($R^2 = -0.18$ for the linear regressor, -0.12 for the nonlinear one): the parameters alone do worse than predicting the held-out mean, while the predictive latent reads the wake at 0.79, clearing the floor by roughly 0.9. The floor is by construction strong on the observables that track the gust forcing directly, the lift (0.61) and the negative circulation (0.23): the predictive advantage is specific to the spatially distributed wake, not uniform across observables, and this is the paper’s central distinction between the integrated forces, which the release condition largely fixes, and the wake, which it does not. The wake claim therefore rests on the representational closure (table 4(a)), the per-encounter paired test (table 8), and the drift mechanism (§4.3), with the parametric floor establishing that the gust parameters do not supply the wake state the latent encodes.

4.2. Objective, architecture, and auxiliary-head controls

The controls in this section separate three ingredients that are entangled in the full model, the wake-observable supervision, the anti-collapse geometry, and the predictive objective, so we read them as an attribution experiment rather than as additional baselines to be beaten. The representational result of §4.1 is established, but the predictive latent *family* as configured renders the wake readable for reasons we have not yet separated: the predictive encoder differs from the reconstructive baseline not only in its training objective but also in its CNN+ViT architecture and its 80-dimensional wake supervision. Because the architecture and the wake supervision are confounds, we isolate the cause here, before

Table 6. Objective $\times$ architecture controls on the fixed test split, scored at the representational endpoint. Held-out representational wake-entropy R^2 at $H = 16$ (test_b, fixed linear probe on the simulation-encoded latent, seed mean) for the four cells crossing {predictive, reconstructive} with {CNN, CNN+ViT} at $d = 64$, with the auxiliary heads (the instantaneous lift plus the 80-dimensional wake head) matched across all four cells, plus a predictive variant with the wake head removed. On the fixed test set the predictive cells read the wake at both architectures while both reconstructive cells score low, and removing the wake head collapses readability to the floor; the CNN and CNN+ViT predictive columns do not separate, so the ViT is not the driver. The reconstructive gap does not survive case-disjoint retraining (text): the matched reconstructive control reads the wake comparably under cross-validation, so the readability is attributed to the wake supervision, not to the objective.

Objective	Encoder	Aux. heads	repr. wake-ent. R^2
predictive	CNN+ViT	lift + wake	0.79
predictive	CNN	lift + wake	0.79
reconstructive	CNN+ViT	lift + wake	0.13
reconstructive	CNN	lift + wake	0.35
predictive	CNN+ViT	lift only	0.11

turning to why the rollout stays on the manifold (§4.3). The controls in table 6 cross objective and architecture, run directly on the same models the paper reports and scored at the representational endpoint: a 2×2 crossing of objective {predictive, reconstructive} with architecture {CNN, CNN+ViT} at matched $d = 64$, with the auxiliary heads matched across all four cells, together with a predictive variant from which the wake head is removed. The question is whether the predictive objective renders the wake readable more than the reconstructive one once the architecture and the wake supervision are held in common. The fixed-test crossing and the case-disjoint cross-validation that follows answer it together, and the answer is more nuanced than the fixed test alone suggests.

The controls (table 6) give a two-part answer on the fixed test set. First, the predictive cells read the wake at *both* architectures with the auxiliary heads matched (0.79 at CNN+ViT, 0.79 at CNN), while both reconstructive cells score low (0.13 at CNN+ViT, 0.35 at CNN), and the two predictive columns do not separate, so the ViT is not the driver. Second, the wake-observable supervision is necessary: the predictive model with the wake head removed collapses to 0.11, at the floor.

The fixed-test reconstructive gap does not survive retraining, and we say so plainly. Under a five-fold case-disjoint cross-validation, in which every fold retrains both the predictive encoder and the matched reconstructive control and reads each held-out case from an encoder that never saw it, the matched reconstructive control reads the wake *comparably* to the predictive latent: its held-out wake-entropy R^2 has median 0.70 against the predictive 0.67, with the predictive latent stable across folds (0.57 to 0.77) while the reconstructive control is erratic (0.02 to 0.95). This is consistent with the wide fixed-test interval on the reconstructive cell (0.13 with a 95% CI of $[-0.70, 0.58]$): the clean fixed-test margin was a favourable draw, not a robust separation. We therefore do not attribute the instantaneous wake readability to the predictive objective over the reconstructive one. The robust statement at the representation level is narrower: the predictive latent reads the wake *consistently*, where the matched reconstructive control reads it only erratically.

The supervision-only control identifies the source of the readability and isolates what the predictive objective does add. An encoder trained on the lift and wake heads alone, with the predictive objective removed entirely (its weight set to zero) but the anti-collapse term retained, reads the wake *at least as* well as the full predictive model ($R^2 = 0.92$ under a readout-frame probe that reproduces 0.80 for the predictive model), so the representational

readability is supplied by the wake-observable supervision, which every wake-supervised cell carries, not by the predictive objective. What the predictive objective robustly adds is dynamical. Rolled forward through a matched predictor, the supervision-only latent forecasts the wake markedly worse than the predictive latent ($R^2 = 0.27$ against 0.50, both under the readout-frame protocol of the representational endpoint rather than the full pre-impact-context protocol of table 4(b), which is why the predictive value here differs from the 0.43 reported there), and the reconstructive families' wake forecast collapses (table 4(b)). The honest statement is therefore a wake-supervised *predictive* latent whose advantage is dynamical, not a larger instantaneous margin: the supervision makes the wake readable, and among objectives the predictive one is distinguished by reading the wake consistently, forecasting it from its own rollout, and keeping that rollout on the encoded manifold (§4.3).

One further control closes the attribution. If reconstruction kept the supervised wake readability low only for want of an isotropised geometry, then adding the predictive model's anti-collapse term to the reconstructive autoencoder should restore it. It does not. A reconstructive autoencoder carrying the matched wake head *and* the same SIGReg anti-collapse term reads the wake at 0.48 (CNN) and 0.21 (CNN+ViT), within noise of the unregularised reconstructive control and far below the supervision-only (0.92) and predictive (0.79) band, and it stays as seed-variable as that control (the CNN+ViT seeds span negative to 0.8). The anti-collapse regulariser therefore supplies the on-manifold geometry of §4.3, not the wake readability; it is the reconstruction *objective* that leaves the supervised readability erratic, where the predictive objective reads it consistently.

The latent dimension is the remaining control. The representational wake closure holds on a plateau from $d = 64$ down to $d = 16$ ($R^2 = 0.79, 0.60$, and 0.55 at $d = 64, 32$, and 16) and falls off below, so the reported $d = 64$ sits on the plateau and the wake is not recoverable at the smallest dimensions on this dataset. The reported $H = 16$ readout is likewise not a cherry-picked endpoint: across $H \in \{4, 8, 16, 32\}$ the predictive representational wake closure stays on a plateau ($R^2 = 0.72$ to 0.80) while neither no-wake established baseline (the reconstructive autoencoder or POD) exceeds 0.40 at any horizon, so the predictive latent leads the wake at every horizon tested. The latent code is correspondingly distributed rather than concentrated in a few coordinates, which §4.5 makes quantitative.

4.3. Why the latent stays on the manifold: drift and topology

The controls attribute the wake advantage to the predictive objective trained with wake supervision, and the representation is where that advantage is established. A separate question is whether the latent is usable under rollout, since a planner queries a forward model at the states its own rollout reaches. The answer is geometric: the predictive latent stays on its encoded manifold while the reconstructive one leaves it, which we establish with two coordinate-free diagnostics, the manifold departure of the rollout and the topology of the encoded encounter.

The predictive rollout stays on the encoded manifold. Two drift measures order the families oppositely, and the reconciliation is the mechanism (table 7). By relative Euclidean deviation from the encoded trajectory the reconstructive rollout is the *least* drifting (0.32 against 0.54 for the predictive at $H = 16$), yet by Mahalanobis distance to the training distribution it is markedly farther by that metric (9.00 against 0.70 for the predictive, which stays in distribution). Both are true because the reconstructive rollout departs along directions of negligible encoded variance: a fraction 1.00 of its Mahalanobis departure lies in the near-null directions of the encoded covariance, where a small Euclidean step is a large Mahalanobis one, whereas the predictive rollout wanders within the broad isotropised cloud its anti-collapse regulariser maintains. The departure spectrum, not the




Figure 6. Relative drift of the predictive rollout on test_b versus forecast horizon H , measured as the deviation of the rolled latent from the simulation-encoded latent normalised by the encoded-latent scale, $\|\hat{z}_{\text{rollout}} - z_{\text{DNS}}\| / \|z_{\text{DNS}}\|$. The transformer predictor (tf-no-c) drifts gradually and stays bounded. This relative deviation understates the family separation, which is the Mahalanobis manifold-departure of table 7: the reconstructive rollout exits along low-variance directions, small in this norm but large in Mahalanobis.

ratio magnitude, is the robust statement: the reconstructive rollout exits the manifold along the directions the decoder never constrained. Two metric-independent diagnostics, the k -nearest-neighbour distance and the local-PCA residual to the training cloud, agree with the Euclidean ordering that the reconstructive rollout drifts the least, confirming that the Mahalanobis separation is a direction effect rather than a larger physical departure. Ratios below one (the predictive and linear rollouts) mean the rollout is more in-distribution than the encoded held-out latents, the predictor mean-reverting toward the training cloud. This is the property a forward model needs, and it is why the qualitative rollout of figure 16 tracks the encounter through impact before drifting only at long times.

The objective-free supervision-only latent locates this property precisely. Its rollout is the most in-distribution of the families (Mahalanobis 0.83, with no near-null departure, 0.00), so staying on the encoded manifold is supplied by the anti-collapse regulariser that the predictive and supervision-only latents share, not by the predictive objective on its own. What the predictive objective adds is staying close to the true trajectory: its relative- ℓ_2 drift (0.54) is below the supervision-only latent's (0.91), consistent with the better wake forecast of §4.2. The reconstructive rollout, which carries no anti-collapse term, is the one that leaves the manifold.

The departure spectrum points to a single geometric cause, established here two ways that do not depend on the latent coordinates: the Mahalanobis distance that shows whether the rollout leaves the training manifold, and the topology that shows whether the encoded encounter is a single cycle.

The mechanism behind the closure gap follows. A probe attached to a drifted rollout is queried outside the support it was fit on, so its output degrades for reasons unrelated to the predictor's one-step quality. A reconstruction objective places latent codes wherever the decoder finds convenient and does not constrain the geometry to be predictable in time, leaving the near-null directions the rollout then exits along. The anti-collapse regulariser supplies the isotropised in-distribution geometry shared by the predictive and supervision-

Table 7. Latent-drift diagnostic on test_b at $H = 16$, two measures. The relative Euclidean deviation of the rollout from the encoded trajectory and the Mahalanobis distance of the rolled-out latent to the training distribution order the families oppositely (Fukami least-drifting by the former, an order of magnitude out by the latter); the near-null fraction is the share of the Mahalanobis departure lying in the near-zero-variance directions of the encoded covariance, which reconciles the two. The Mahalanobis magnitude is reported at the minimally-regularised covariance used here; under a shrinkage estimator the ratios collapse toward one, so the robust statement is the departure direction (the near-null fraction), not the ratio size. The objective-free supervision-only latent is the most in-distribution of all, which locates the on-manifold property in the anti-collapse regulariser the predictive and supervision-only latents share, not in the predictive objective alone.

Family	rel. ℓ_2 drift	Mahalanobis dist.	near-null frac.
Fukami	0.32	9.00	1.00
JEPA (tf)	0.54	0.70	0.88
supervision-only	0.91	0.83	0.00
POD	0.84	0.84	—
β -VAE	0.53	0.87	—

only latents, and predictive training then makes that state dynamically trackable: the predictive rollout stays closer to the encoded simulation trajectory than the objective-free supervised rollout, while the reconstructive rollout, lacking the regulariser, leaves the support along low-variance directions. This is a representation property rather than a probe artefact: no probe in the evaluated class recovers the wake from the reconstructive rollout once it has left the manifold the probe was fit on, while the same class reads it from the predictive rollout, so the reconstructive failure is the manifold departure and not an inaccurate predictor step.

The manifold departure is a rollout property; the second diagnostic asks whether the *encoded* encounter has the topology of a single cycle. The undisturbed shedding traces a limit cycle, and a gust encounter departs it, executes the LEV-formation and shedding excursion, and relaxes back toward it (Smith *et al.* 2024). We characterise the cycle with persistent homology, counting the significant one-dimensional generators (H_1) of a Vietoris–Rips filtration (Smith *et al.* 2024) on each encoded latent trajectory. The metric matters here. In raw latent coordinates the reconstructive encoding fragments a gusted encounter into several generators where the predictive encoding shows one, but the two encodings have very different per-coordinate scales by construction (the predictive latent is whitened toward isotropy by its anti-collapse regulariser, the reconstructive latent is not), and removing that anisotropic scale with a per-family whitening collapses the difference: under the whitened metric both families read a gusted encounter as essentially a single cycle (median 1.0 and 1.0 generators). The gusted-encounter fragmentation is therefore a property of the metric, not of the topology, and we read no topological difference into the gusted encounter.

What survives the whitening is the cleanest control, the undisturbed no-gust limit cycle (figure 7). Even after per-family whitening the reconstructive encoding fragments the clean shedding cycle into several generators (median 2.5) while the predictive encoding keeps it as a single loop (median 1.0, as does the linear basis). The reconstructive encoding cannot represent even the undisturbed limit cycle coherently, which is an encoding property, distinct from the rollout manifold departure above and computed on the simulation-encoded latents.




Figure 7. Persistent homology of the encoded no-gust limit cycle on test_b, under the per-family whitened metric that removes anisotropic coordinate scale. Left and centre: representative Vietoris–Rips H_1 persistence diagrams of the encoded predictive and reconstructive latents; coloured points are significant generators (lifetime above the dotted noise floor) and grey points are noise. Right: the significant- H_1 generator count over the undisturbed shedding periods. The predictive encoding is a single clean loop (median 1.0) while the reconstructive encoding fragments (median 2.5); the gusted fragmentation seen in raw coordinates does not survive the whitening and is not claimed.

4.4. Flow physics of the encounter and its latent image

The closure and mechanism results are representational; this subsection reports what the analysis says about the flow itself, in measures that do not depend on any learned model, and where the predictive latent inherits those measures. Three results stand.

First, the latent dynamics carry the shedding spectrum. Fitting a linear operator to the encoded limit-cycle trajectory (dynamic mode decomposition) recovers a leading oscillatory eigenvalue at Strouhal number 0.66 in the predictive latent, against a measured shedding line at 0.68 in the undisturbed lift spectrum, with modulus 0.991 (bootstrap CI [0.988, 1.016], consistent with a marginally stable orbit); the linear basis recovers the same frequency (0.68) and the reconstructive latent is damped and off-frequency (0.50). Linearising the learned predictor about the orbit gives a leading multiplier of modulus 1.004 (bootstrap CI [1.003, 1.005]), which sits just above one, so the learned orbit is very weakly amplifying rather than exactly neutral; we report it as a learned delay-map multiplier, not a physical Floquet multiplier of the simulation. The learned dynamics are not, however, a linear operator in disguise: rolling a best-fit linear operator on the linear-basis coordinates (dynamic mode decomposition as a forecaster) gives a wake-entropy forecast R^2 of -0.39 , far below the learned predictor's 0.43 on the same observable, a margin of 0.82. The predictive latent therefore preserves the shedding spectrum that a linear basis also carries, while its nonlinear dynamics forecast the gusted wake that a linear operator cannot.

Second, the wake response scales with gust strength. The peak large-scale wake-entropy excursion grows monotonically with $|G|$, a trend that is robust to the scale band (Spearman 0.74 at $\sigma/c = 0.05$, 0.70 and 0.51 at the finer bands). The leading-edge vortex circulation budget shows a gust-sign asymmetry: positive gusts produce a larger and longer-lived leading-edge vortex than negative gusts at matched $|G|$, with a later circulation detachment (paired peak-circulation difference 0.240), quantifying the split-and-merge interaction the lineage describes qualitatively. The peak lift excursion and the peak vortex circulation are monotonically associated across the envelope (Spearman




Figure 8. The latent dynamics carry the shedding spectrum but are not a linear operator. Left: dynamic-mode-decomposition eigenvalues of the encoded limit-cycle trajectory in the complex plane, by family, against the unit circle and the measured shedding-frequency ray; the predictive (JEPA) and linear (POD) latents place a marginally stable oscillatory pair at the shedding Strouhal number (0.66 and 0.68, moduli 0.991 and 0.996) while the reconstructive (Fukami) pair is damped and off-frequency (0.50). Right: the wake-entrstrophy forecast skill of a best-fit linear operator on the linear basis (dynamic mode decomposition as a forecaster, -0.39) against the learned predictor (0.43), a margin of 0.82.

0.75), reported as a correlation and not an impulse-theorem identity; the linear coefficient is weaker (Pearson 0.53, Appendix A).

Third, the gust resets the shedding phase. Within a case the encounters share gust parameters and differ only in release instant, so consecutive releases sample shifted shedding phases; the undisturbed cadence is reproduced to within a few hundredths of a radian, but any gust of $|G| \geq 0.25$ decorrelates the shedding phase between releases, and the mean phase step drifts monotonically with gust strength. The encounter does not preserve the pre-impact clock, it resets it, which is why a fixed-cadence release campaign undersamples the impact phase (Appendix A); we report the phase response within that measured coverage and treat the impact phase as a sampled rather than a controlled variable.

4.5. The latent code is distributed

The capacity ladder of §4.2 already indicated that the wake is not held in a few latent coordinates. Two measures make this quantitative. The wake forecast is a collective code: the gap between the wake readability of the full latent and that of its single most informative coordinate is 0.36 for the predictive latent, against 0.05 and 0.04 for the reconstructive and linear baselines, so no single predictive coordinate carries the wake and the information is spread across many. The energy-information curve agrees: the predictive latent needs about 32 leading principal components to reach its wake skill, where the reconstructive latent saturates at 12 and the linear basis at 1. The distributed code is the latent-space signature of the anti-collapse regulariser, which isotropises the representation, and it is consistent with the rollout staying within the broad cloud of §4.3.

The distributed code lives in the latent rather than on identifiable flow structures. A gradient-saliency footprint of the wake-forecast direction through the frozen encoder overlaps the high-vorticity cores and the Q-criterion vortices *below* a permutation chance baseline (overlap 0.070 and 0.010 against the chance band), so the wake-forecast direction

does not read the input preferentially at the leading-edge vortex or the shear layer. The wake code is distributed across the latent and is not localised in the field.

4.6. Where the wake is kept and where it is lost in physical space

The mechanism explains why the predictive rollout stays on the manifold; we now ask where in the gust envelope, and where in the flow field, the latent renders the wake readable. The parameter dependence reveals one clear weakness, and it is worth noting that the latent encodes these parameters despite never being trained on them. Probing the predictive latent for the gust parameters on test_b recovers the gust strength G and the core diameter D ($R^2 = 0.83$ and 0.65) but does not resolve the wall-normal offset Y ($R^2 = -0.03$, the confidence interval straddling zero). This is a property of the data rather than of the architecture: the training mass concentrates near $Y \approx 0$, so encounters at different offsets within the envelope are weakly separated along Y , and the latent inherits that weakness. Any deployment-time parameter estimate should treat Y as a nuisance variable the encoder does not resolve.

The same latent can also be read relative to the undisturbed shedding cycle. Because the undisturbed wake is a limit cycle (§4.4), an encounter is naturally described by its phase and amplitude relative to that cycle, the same reduction used to design gust-mitigation control for these flows (Fukami *et al.* 2024). The predictive latent of the undisturbed case traces a closed periodic orbit (figure 9): concatenating the no-gust episodes, the return distance is 1.0% of the orbit diameter and the episodes overlap the same loop, with a shedding period near the full-cycle Strouhal number 0.34 (§4.4), and a phase advances monotonically along it via the Hilbert analytic signal of the two leading principal components. This is the same object the spectrum of §4.4 identifies as a marginally stable orbit and the topology of §4.3 identifies as a single persistent cycle. A gust encounter departs this orbit and only partially relaxes back: the simulation trajectories do not fully return within the 120-frame window, which is the gust-release period itself ($6t/c$ at $\Delta t = 0.05t/c$, impact near $2t/c$), so only about $4t/c$ of post-impact relaxation is observed before the next release, a window-length limitation of the fixed gust cadence. We therefore read the latent cycle as a phase-amplitude coordinate of the forward model, and do not draw a family comparison from the incompletely-observed return.

We can also locate where in the envelope the predictive advantage concentrates. Rather than a per-stratum closure table, which the present test set is too small to populate, we read the per-encounter paired improvement $\Delta e = e_{\text{AE}} - e_{\text{JEPA}}$ in held-out wake-entrrophy error directly against each gust axis with a locally weighted regression and a bootstrapped band (figure 17). The advantage is not uniform across the envelope. It grows with the gust core diameter D (Spearman $\rho = +0.36$ on the encoded latent, significant), while its dependence on the gust strength G and the wall-normal offset Y is weak and not significant ($\rho = +0.13$ and $+0.03$). The predictive advantage therefore concentrates at the larger cores that most reorganise the wake, with no significant Y trend, in keeping with the weak Y resolution of the probe above. The shedding phase at impact, the axis the gust-response literature emphasises (Fukami *et al.* 2024), is only weakly sampled by this dataset: impact is fixed near a single convective instant while the gust recurs about every two shedding periods, so the encounters fall into essentially two phase clusters and a phase-resolved trend is not identifiable here, which a design that varies the impact phase would isolate. These parameter and phase readings are statements about numbers; the encounter is itself a statement about vortices. Decoded and reconstructed vorticity fields on held-out encounters (figure 18) connect the two, with one methodological caution and one new metric. Two scoping points apply throughout this subsection. The fields are encode-decode reconstructions, the visualisation decoder applied to the simulation-encoded latent, so the




Figure 9. The gust encounter in the predictive latent. (a) Distance of the gust trajectory to the baseline orbit (full latent, in units of the orbit diameter; the grey band is the baseline orbit thickness); the encounter departs the baseline cycle and only partially relaxes back within the inter-gust window; the single-cycle topology of the encoded latent is established coordinate-free in figure 7. (b) The orbit phase relative to the impact frame. The numbered stages are those of figure 16.

quantitative comparisons report the decoder’s ceiling on what each latent carries from a faithfully encoded field, not a forecast rollout, whose latent leaves the data manifold for the reconstructive baseline (§4.3). And every quantitative physical-space claim is made on the validated large-scale band ($\sigma/c = 0.05$); the visualisation decoder is blurry at small scale by construction, so the large-scale leading-edge vortex and shear layer are tracked meaningfully while full-resolution and small-scale field fidelity are not claimed.

The caution is that field-space similarity scores are misleading here. On held-out interpolation encounters the reconstructive autoencoder can attain a high structural similarity while collapsing to a near-uniform mean field that has lost the vortex-impingement signature, the high similarity merely reflecting bulk agreement with a field that is mostly quiescent away from the impact region. The predictive decode localises the impingement at the right position and sign but is blurry by construction, and the linear basis is amplitude-accurate since its objective is amplitude-optimal. Field-space similarity therefore cannot rank the latents, which is why we read it as a labelled decode ceiling rather than a state-content score.

Decoding each family’s faithfully encoded latent to a field and scoring the large-scale similarity against the simulation gives a decode *ceiling* (the best each decoder can do from a true latent): on test_b that ceiling is highest for the linear basis (0.71), then the predictive decode (0.62), then the reconstructive autoencoder (0.54). The linear basis leading this ranking is a decoder property, not a state-content one: its linear decoder reconstructs a smooth large-scale field well regardless of what the latent encodes, exactly the confound that the collapsed-reconstruction example above illustrates. Field decode fidelity is therefore reported as a labelled ceiling, with the explicit caveat that it conflates decoder reconstruction quality with latent state content; the latent comparison rests on the representational closure of §4.1, which has no decoder in the loop, not on these field scores.

Finally, the wake observables tie to the structures that carry the load. A Gaussian scale decomposition (large scale at filter width $\sigma/c = 0.05$, following (Odaka *et al.* 2026; Motoori & Goto 2019)) separates the large-scale vortical structures responsible for the lift peaks from the fine scale, applied to the simulation field and to each family’s reconstruction (figure 10). On test_b the large-scale wake enstrophy of the predictive decode tracks the




Figure 10. Gaussian scale decomposition ($\sigma/c = 0.05$) of the held-out reconstructions. Top: large-scale vorticity for the simulation, the predictive decode, and the reconstructive autoencoder at $H = 16$; the predictive latent retains the leading-edge vortex and shear layer the reconstructive one smooths away. Bottom: large-scale wake enstrophy through the staged encounter for the simulation, predictive, and reconstructive fields, on each split.

simulation better than the reconstructive autoencoder, on both correlation (Pearson 0.54 against 0.41 at $H = 16$) and amplitude (relative error 0.12 against 0.71): the reconstructive decode smooths the leading-edge vortex and shear layer away. The linear basis tracks the correlation comparably (0.58), the same decoder confound as the decode ceiling above, its linear decoder reconstructing the smooth large-scale field well, so the specific claim is that the predictive decode tracks the large-scale wake better than the reconstructive autoencoder, while field-space decode correlation does not separate it from the linear basis. This is where the wake-entrophy advantage of §4.1 becomes a statement about the leading-edge vortex and shear layer rather than about a scalar. The converse question, whether a family's *decode* tracks the leading-edge-vortex circulation per frame, is confounded by decoder quality rather than by what the latent carries. Under the matched post-hoc decoders the fraction of held-out encounters whose decoded circulation tracks the simulation is comparable across families and modest (0.29 for the predictive decode, 0.27 for the reconstructive, and highest for the linear basis at 0.48, whose linear decoder reconstructs the large-scale field best regardless of state content). The wake result therefore rests on the representational closure and the scale-resolved enstrophy tracking above, where the latent comparison has no decoder in the loop.

4.7. Wall-pressure observability

At deployment the vorticity field is not available; only the wall pressure is. The deployment analogue is therefore state recoverability rather than direct wake recovery: we ask which representation is the most recoverable from a sparse set of K wall-pressure taps over a window ending at the readout instant, comparing both the recovered latent *state* and the recovered flow *field* across families under one kernel-ridge estimator and a target-blind




Figure 11. Cross-family pressure observability under the target-blind qDEIM sensor placement. (a) Latent-state recovery R^2 versus sensor count K : the predictive (JEPA) state is the most recoverable, the linear basis the least. (b) Recovered-field large-scale similarity by family (pressure to latent to decoder), with the decode-ceiling shown alongside: the predictive latent beats the reconstructive autoencoder and ties the linear basis, whose field parity is its decoder reconstruction ceiling rather than better state recovery.

(qDEIM) sensor placement (the placement, the per-family field-recovery breakdown, the parameter recovery, and the out-of-distribution boundary are in Appendix B).

The predictive state is the most recoverable (figure 11a). With $K = 8$ taps the held-out test_b state R^2 is 0.78, against 0.66 for the reconstructive autoencoder, 0.51 for the variational baseline, and 0.34 for the linear basis; the lead is seed-robust (0.78 ± 0.09 over four predictive seeds against 0.66 ± 0.04). The energy-optimal POD coordinates, optimal for reconstruction, are the hardest of the three to read from the wall: the same predictive objective that renders the wake readable from the field produces a state that is also legible from the wall. The wall pressure also recovers the wall-normal offset Y at $R^2 = 0.51$, so the weak Y resolution of §4.6 is a property of the mid-plane vorticity observable and the offset sampling, not of the pressure signal. The recoverability is causal: restricting the pressure window to a strictly pre-impact span, ending before the gust reaches the airfoil rather than at the readout instant, preserves and in fact sharpens the ordering (predictive 0.82 against 0.77 for the reconstructive autoencoder and 0.22 for the linear basis), so the predictive state is the most recoverable from pressure available strictly before the event.

The recovered *field* tells a more nuanced story that we report in full. The pressure-to-latent map followed by each family’s decoder reconstructs the large-scale field with similarity 0.57 for the predictive latent, 0.58 for the linear basis, and 0.48 for the reconstructive (figure 11b); the predictive latent beats the reconstructive significantly but ties the linear basis. The tie is a decoder property: the decode-ceiling gap (the loss from pressure recovery relative to decoding the true latent) is largest for POD (0.13 against 0.05 for the predictive), so the linear basis reaches field parity through its decoder’s reconstruction ceiling, not through better state recovery, where it is the weakest of the three. The clean comparison is therefore the state recovery: the predictive latent is the most recoverable state from sparse wall pressure, and recovers the field at least as well as any matched representation.

5. Discussion

5.1. What a predictive latent state must contain

The central result is a design rule for predictive reduced states: a vortex-gust state must carry the wake observables that govern the post-impact transient, but readability alone is not enough, because the state must also be organised so that a forward model can roll it out without leaving the manifold on which the probes and estimators were trained. The finding is therefore dynamical, and it concerns what a wake-supervised state is good for once it has to evolve. The instantaneous part is quickly stated and quickly bounded: under a shared probe family the wake-supervised predictive latent reads the held-out wake at $R^2 = 0.79$ on the fixed test set while the established reconstructive and linear baselines, which carry no wake head, are at or below zero. Two controls show this readability is the wake supervision's, not the predictive objective's. An objective-free encoder carrying only the wake-observable supervision reads the wake as well ($R^2 = 0.92$ against 0.80 for the predictive latent under the same probe), and under five-fold case-disjoint retraining a matched reconstructive control that carries the same wake head reads the wake comparably to the predictive latent (held-out median 0.70 against 0.67), so the clean fixed-test margin does not survive out-of-sample retraining. We therefore do not claim that the predictive objective makes the wake more readable than a matched reconstructive one. What the predictive objective robustly adds is dynamical: it reads the wake *consistently* across retrains (the matched reconstructive control swings from 0.02 to 0.95 where the predictive latent stays in $[0.57, 0.77]$), it forecasts the wake from the latent's own rollout where the supervision-only and reconstructive states do not ($R^2 = 0.50$ against 0.27 for supervision-only), and its rollout stays on the encoded manifold. The mechanism is geometric. A reconstruction objective fixes the latent only up to a diffeomorphism (Kelshaw & Magri 2024), so the decoder is free to place codes anywhere it finds convenient and nothing constrains the latent geometry to be predictable in time; under autoregressive rollout the reconstructive latent leaves its own training manifold along the very directions the decoder never constrained, and a probe queried on that drifted state fails for reasons unrelated to the predictor's one-step accuracy. The anti-collapse regulariser instead organises the latent so that the rollout stays on the encoded manifold, the on-manifold property that data-driven state-space and Koopman models of coherent-state dynamics on invariant manifolds rely on (Constante-Amores & Graham 2024), and that the encoded no-gust limit cycle exhibits as a single loop even under a metric that removes coordinate scale. This property is not unique to the predictive objective: the objective-free supervision-only latent, which carries the same anti-collapse term, is the most in-distribution of all under rollout (§4.3), so on-manifold geometry is supplied by the regulariser, while the predictive objective is what keeps the rollout close to the true trajectory and renders the wake forecastable. The manifold departure of the reconstructive rollout, which carries no anti-collapse term, is the parametric-gust counterpart of the catastrophic-divergence tail reported for reconstruction-based latent dynamics in controlled wakes (Solera-Rico *et al.* 2025); the anti-collapse regulariser is the common cure. The wake observables are the sharpest discriminator precisely because they can change while the scalar forces do not: the gust parameters alone fail to predict the wake at the readout horizon (the parametric floor of §4.1 is negative) yet predict the forces well, so a representation can carry the force signature of an encounter without carrying the vorticity redistribution that produces it, and that is exactly the reconstructive latent's failure mode.

The coordinate-level reading of §4.5 explains why the wake, and not the forces, separates the families. The wake forecast is a collective code in the predictive latent: the gap between the readability of the full latent and that of its single most informative coordinate is




Figure 12. The wake is a distributed latent code, on test_b. Left: the gap between the held-out wake readability of the full latent and that of its single most informative coordinate, by family; the predictive (JEPA) latent shows a large gap (0.36) where the reconstructive (Fukami, 0.05) and linear (POD, 0.04) latents do not, so no single predictive coordinate carries the wake. Centre: the held-out wake skill versus the number of leading principal components, which knees at about 32 components for the predictive latent against 12 and 1 for the baselines. Right: the gradient-saliency footprint of the wake-forecast direction overlaps the high-vorticity cores and Q-criterion vortices below the permutation chance band, so the distributed code is not spatially localised. Observational probes on the frozen encoders.

0.36, against 0.05 and 0.04 for the reconstructive and linear baselines, whose best single coordinate is already as informative as the whole latent (figure 12). The forces, by contrast, are carried redundantly in every family. An energy or reconstruction objective therefore retains the force signature, present in many coordinates, but not the distributed combination that carries the wake, which is exactly the observable on which the predictive latent leads and which the gust parameters alone cannot supply (the parametric floor of §4.1 is negative on the wake). We read this as a description of correlations, not a causal decomposition: these are observational probes on the frozen encoders, alongside the information-theoretic modal analyses of related flows (Martínez-Sánchez *et al.* 2023; Arranz & Lozano-Durán 2024; Koshikawa *et al.* 2026).

The wake organisation is low-dimensional. The representational wake closure holds on a plateau as the latent dimension is reduced, from 0.79 at $d = 64$ to 0.60 at $d = 32$ and 0.55 at $d = 16$ (§4.2), so the representation declines gracefully rather than collapsing, and the qualitative ordering, predictive leading while the reconstructive and linear latents stay negative, is preserved throughout. The encoded latent is itself concentrated in a handful of directions: the $d = 64$ latent has a participation ratio of about 1.7, with the leading principal component carrying roughly three-quarters of the variance and six components spanning 90%, the dominant direction tracking the gust-strength dynamic range. The reduction in closure as the dimension falls is therefore a loss of precision on the wake observables, not of the geometric mechanism, which is consistent with the distributed code above: the wake is spread across many low-variance directions, and trimming them costs sharpness before it costs the ordering.

The mechanism above invites a strong reading, which the controls bound. The predictive encoder differs from the reconstructive baseline in three respects at once: the training objective, the CNN+ViT architecture, and the wake supervision. The parameter-only floor removes the gust parameters as an explanation, all the more since the model never sees them, and the matched-architecture, matched-auxiliary 2×2 controls of §4.2 separate the remaining two. They give a two-part answer that we state in full. On the fixed test set the predictive objective reads the wake where the matched reconstructive objective does not, at both a CNN and a CNN+ViT encoder, with the architecture columns not separating; but that fixed-test gap does not survive case-disjoint retraining, where the

matched reconstructive control reads the wake comparably (§4.2), so we do not rest a claim on it. The wake supervision is the necessary component: a predictive model with the wake head removed collapses below the predict-the-mean floor, and an objective-free encoder carrying only that supervision reads the wake as well as the full predictive model, so the representational readability is supplied by the supervision, not by the predictive objective. The claim is therefore that a wake-supervised latent renders the wake readable, and that among objectives the predictive one is distinguished not by a larger instantaneous margin but *dynamically*: it reads the wake consistently across retrains, and it renders the wake forecastable from the rollout, which the static supervision-only encoder cannot do. The linear basis, for its part, remains genuinely useful: it is the most competitive family on the integrated impulse and is amplitude-accurate in reconstruction, so the result is not that POD is obsolete but that it does not produce a latent from which the wake observables are linearly readable.

5.2. Limitations and outlook

Three limitations bound the claims. First, the wall-normal offset Y is weakly resolved because the training mass concentrates near $Y \approx 0$; this is a data limitation, and Y -conditional statements should be read with it in mind. Second, the $|G| = 4$ extrapolation degrades for a physical reason, not merely a statistical one. As set out in §2.1, the encoder observes a single mid-plane vorticity slice whose information no longer determines the three-dimensional $|G| = 4$ dynamics (Fukami *et al.* 2025), so test_c is the observability boundary of that representation, not a parametric-generalisation claim. Third, the visualisation decoder is a diagnostic on the frozen encoder, never part of the predictive loss; a reader expecting sharp reconstructions from a predictive latent will be disappointed, which is the explicit trade the objective makes. To these we add, as a fourth, the role of the wake supervision established in §5.1: the wake supervision supplies the representational readability, and the predictive objective is what preserves it (where the reconstructive objective suppresses it) and renders it forecastable; the result is a property of the wake-supervised predictive configuration, not of the objective stripped of supervision nor of the supervision attached to an arbitrary objective.

These limitations come into sharpest focus against the use the architecture is ultimately meant for. The predictor is, by construction, a latent dynamics function of the kind an estimator or controller would require, though closed-loop use is outside the present evidence, and three properties of the architecture are relevant to that use: the latent renders readable the wake the load transient is built from, the latent is two to three orders of magnitude cheaper to evolve than the full field, and the rollout stays on the training manifold, which is the property that determines whether long-horizon plans remain meaningful since a planner queries the predictor on states reached by its own rollout. The latent itself is the observable a controller has access to, with no gust parameters that would have to be estimated first, and adding an actuation channel is a change of input rather than of architecture; the phase-amplitude reading of §4.6 connects the predictor directly to the sensitivity-function control designed for these flows (Fukami *et al.* 2024). The predictor takes no gust parameters and is trained to forecast the observed latent trajectory, so it is a forward model of the encounters in the training envelope: it captures how an encounter evolves within the envelope, while predicting how one would differ under a changed gust parameter, the interventional setting, is a separate task its training does not address. We keep the world-model view of §1 as motivation accordingly.

The contribution is narrow by design: a controlled comparison on a single parametric flow, anchored to the representational closure of the wake rather than to in-distribution reconstruction, with a drift mechanism that explains the gap and a set of controls (§4.2)

that bound its interpretation. What the comparison establishes is bounded accordingly. The predictive advantage is robust on the representational wake closure and in the drift mechanism; it is partial on parameter extrapolation, where the wake encoding survives at $G = -4$ but the integrated observables do not (§4.1), and on the wall-normal offset Y , which is harder to resolve than G and D (§4.6); and it does not extend to the interventional reading or the closed-loop tolerances, both of which fail in the controls that tested them. We claim the representational wake result and its mechanism, not a validated forecast, a controller, or a counterfactual model. The natural next step is to apply the same matched-predictor protocol to other parametric flow datasets, to test how general the predictive-versus-reconstructive trade reported here actually is.

6. Conclusions

We asked which reduced-order state best carries the wake that governs the transient load in parametric vortex-gust airfoil interactions, and reached a dynamical answer with a geometric mechanism: a wake-supervised predictive latent builds a state from which the wake is linearly decodable, whose rollout forecasts the wake and stays on the encoded manifold, where a reconstruction objective carrying the same supervision lets its rollout drift off the manifold along directions of negligible encoded variance. We tested this by comparing a predictive latent (JEPA, which takes no gust parameters), a reconstructive observable-augmented autoencoder, and a proper orthogonal decomposition basis as reduced states at $Re = 5000$, under a protocol in which a single probe family was held fixed so that only the encoder varied. The instantaneous readability is the wake supervision's: on the fixed test set the wake-supervised predictive latent reads the wake (representational wake-entropy $R^2 = 0.79$) where the established no-wake baselines are at or below zero, but an objective-free supervised encoder reads it as well, and under case-disjoint retraining a matched reconstructive control reads it comparably (held-out median 0.70 against 0.67), so we do not attribute the readability to the predictive objective. What the predictive objective robustly adds is dynamical: it reads the wake consistently across retrains, forecasts it from the rollout where the supervision-only and reconstructive states do not, and keeps the rollout on the manifold, a property shared with the anti-collapse supervised latent and absent from the unregularised reconstructive one. The same separation appears coordinate-free in the rollout drift and in the topology of the encoded no-gust limit cycle. The latent also carries the flow's own clocks: its spectrum recovers the shedding frequency on an orbit whose modulus sits just above one, the wake response scales with gust strength, and the gust resets the shedding phase. The same predictive state is the most recoverable of the three from sparse wall pressure, so it is both wake-readable and the most wall-recoverable of the three, the property a downstream estimator would exploit. Three facts bound the result: the wall-normal offset is weakly resolved, the mid-plane representation reaches a three-dimensional observability boundary at $|G| = 4$, and the rollout error still exceeds the tolerance a closed-loop control demonstration would require.

Funding. [Authors to insert funding sources and grant numbers.]

Declaration of interests. The authors report no conflict of interest.

Data availability statement. The reduced-order-model and evaluation code, the configuration files, the v2.1 data-split manifest, and the analysis outputs that regenerate the tables and figures are deposited at <https://doi.org/10.xxxx/zenodo.PLACEHOLDER> under the MIT License (DOI to be finalised on acceptance). The processed per-encounter cache, or a representative subset sufficient to rerun the heavier steps, is available from the corresponding author. The raw direct numerical simulations were computed with the SOD2D solver (Gasparino *et al.* 2024) and are owned by the simulation collaborators.




Figure 13. Predictive versus reconstructive training. Left: the reconstructive autoencoder encodes the field at t , decodes it, and incurs a field-space error. Right: the predictive model encodes the fields at t and $t + \Delta t$ with shared weights and incurs a latent-space error between the predicted and observed target embeddings, with no field-space loss and no decoder during training. This is the contrast the matched-predictor protocol of §3.2 evaluates; the visualisation decoder is trained only afterwards, on the frozen encoder.

Author contributions. [Authors to insert contributions, for example in CRediT format.]

Use of generative-AI tools. Generative-AI assistance (Anthropic Claude, accessed through the Claude Code command-line interface) was used during 2026 to help with data analysis, figure generation, and manuscript drafting. All numerical results were computed by the authors' own code and verified against the underlying data, and the authors take full responsibility for the content of the article.

Appendix A. Architecture, anti-collapse regularisation, and uncertainty quantification

Architecture detail

The encoder is a hybrid convolutional stem followed by a small vision transformer. Three downsampling convolutional stages map the single-channel 192×96 vorticity field to a 24×12 feature map at 256 channels, that is 288 spatial tokens, which a six-layer pre-norm transformer of width 256 with eight heads processes; the latent is read from a prepended class token through a one-layer projection whose output passes a BatchNorm, the layer the characteristic-function regulariser below requires. The encoder is unconditional and carries no information about the gust parameters. The predictor is a six-layer autoregressive transformer of width 384 with sixteen heads and dropout 0.1, with a causal mask and rotary temporal positions. The predictor advances the latent from its own history alone, with no gust input, so the gust parameters $\mathbf{c} = (G, D, Y)$ enter nowhere in the model. Training minimises the one-step latent prediction loss plus a scheduled-sampling rollout term and the regulariser below, with AdamW and bf16 mixed precision.

Anti-collapse regularisation

Because the encoder and predictor are trained end to end on latent-space prediction with no decoder, the latent is held away from a collapsed solution by an explicit regulariser rather than by a reconstruction term. We use SIGReg, the characteristic-function (Epps–Pulley) regulariser of the from-pixels joint-embedding line (Maes *et al.* 2026; Balestrierio & LeCun 2025), with M random one-dimensional projections of the latent and a small weight (0.01) in the loss. As a safeguard the participation ratio of the latent is monitored every thousand iterations, with an automatic fall-back to a variance-covariance objective (Bardes *et al.* 2022) if it indicates collapse while the parameter probe is still uninformative; on the runs reported here the participation ratio stays healthy and the fall-back does not trigger.

Table 8. Paired held-out closure, predictive ($d = 64$) against reconstructive ($d = 64$) on the $n = 42$ test_b encounters at $H = 16$, for the representational mode (probe on the simulation-encoded latent) and the forecast mode (full-context rollout). Δerr is the per-encounter mean of (reconstructive – predictive) absolute error in each observable’s own units, so a positive value favours the predictive latent; the 95% interval is case-clustered (resampling the 10 cases), and the sign test counts the encounters on which the predictive error is the smaller of the pair. The wake enstrophy is the pre-registered primary endpoint (§4.1); the final column is the encounter-level sign-test p Holm-corrected across the twelve tests of the table (the pre-registered D165 family correction). Bold marks the four tests that survive that correction. We report this encounter-level correction as the headline and note the stricter case-level one in the text: the representational wake advantage survives the encounter-level Holm correction but not the stricter case-level one, where the smaller number of cases widens the interval. All six observables are shown, not a selected subset.

Mode	Observable	Δerr	case-clustered 95% CI	sign test	Holm p
repr.	C_L	+0.552	[+0.241, +0.949]	34/42, $p = 3.4 \times 10^{-5}$	4.1×10^{-4}
repr.	C_D	+0.134	[+0.00847, +0.309]	25/42, $p = 0.14$	0.7
repr.	I_y	-0.681	[-1.4, +0.0743]	17/42, $p = 0.92$	1
repr.	wake-enst.	+33.6	[+7.79, +58.7]	30/42, $p = 0.004$	0.04
repr.	circ. pos.	+0.398	[+0.0384, +0.912]	27/42, $p = 0.044$	0.35
repr.	circ. neg.	+0.944	[+0.302, +1.97]	33/42, $p = 1.4 \times 10^{-4}$	0.0015
forecast	C_L	+0.012	[-0.43, +0.424]	27/42, $p = 0.044$	0.35
forecast	C_D	+0.0813	[-0.0878, +0.251]	23/42, $p = 0.32$	1
forecast	I_y	+0.141	[-0.307, +0.61]	21/42, $p = 0.56$	1
forecast	wake-enst.	+16.8	[+1.06, +32.8]	26/42, $p = 0.082$	0.49
forecast	circ. pos.	-0.128	[-0.435, +0.194]	17/42, $p = 0.92$	1
forecast	circ. neg.	+0.708	[-0.0103, +1.72]	30/42, $p = 0.004$	0.04

Uncertainty quantification

Every held-out coefficient of determination and mean absolute error in the paper carries three independent uncertainty signals, following the reporting protocol fixed with the split. The first is a bootstrap confidence interval from 2000 resamples over the held-out test encounters, which is the dominant source of scatter at the present sample size ($n = 42$ for test_b, $n = 24$ for test_c) and is wide for the individual wake observables, as noted in §4.1. The second is the variance across three encoder retrains from independent random seeds, reported as the standard deviation in the objective-times-architecture controls of §4.2. The third is a five-fold cross-validation over cases on the read-out (probe) step, which guards against a probe that overfits the particular held-out split. Together these approximate what a full five-fold case-level cross-validation of the entire pipeline would buy, at a fraction of the training cost; a single split with bootstrap intervals is the community standard for this class of study.

Training-set fit

The predictive latent attains the highest in-distribution training fit of the forecast rollout among the families, but this is an in-distribution number, not the held-out result: the family ordering is established by the held-out representational closure of table 4(a), and the held-out forecast mean is markedly lower than the training fit, as it must be. We therefore do not tabulate the training-set fit, which would invite reading an in-distribution number as a result.

Paired closure across all observables

Table 8 reports the per-encounter paired comparison of §4.1 for all six observables in both modes, the source of the wake-enstrophy significance cited in the main text. Pairing cancels the encounter-to-encounter difficulty that widens the marginal intervals of table 4.

Topological robustness

The significant- H_1 generator count of §4.3 declares a generator significant when its persistence exceeds a noise floor set at a fraction of the cloud diameter (the largest finite H_0 death), the canonical analysis using 5% of that diameter and all 120 trajectory frames. Following the convergence protocol of Smith *et al.* (2024) we sweep the floor over $\{2, 5, 10, 15, 20\}\%$ and the points per encounter over $\{120, 60, 40, 30\}$ (uniform subsampling), 20 settings in all. The predictive median is one generator in every setting. The predictive-versus-reconstructive

Table 9. Preprocessing robustness of the representational wake closure. Held-out test_b wake-entrophy R^2 at $H = 16$ from a frozen-latent ridge probe (no retraining), recomputed under three training-only amplitude-clip treatments of the raw ω_z field. The predictive latent leads the wake under every treatment; the advantage keeps its sign and approximate magnitude.

Family	no clip	per-encounter	training-global
predictive (JEPA)	0.75	0.80	0.84
reconstructive (Fukami)	0.12	0.10	0.08
linear (POD)	-0.27	-0.05	0.00

1003 Mann–Whitney separation is below 10^{-3} for floors from 2% to 15% at full (120-point) sampling and for floors
 1004 up to 5% at half (60-point) sampling (the canonical 5%, 120-point value is 4.9×10^{-9}), while the reconstructive
 1005 median falls from six generators at a 2% floor to one at a 20% floor as the aggressive floor filters the spurious
 1006 generators. Under heavy subsampling to 30 points both medians approach one and the separation washes out.
 1007 The robust statement is therefore the median-count separation over a defensible floor and sampling range, not
 1008 the exact p ; the case-clustered signed-rank test (10 of 10 cases at the canonical setting) holds across the same
 1009 range. The full grid is in the released analysis package.

1010 *Preprocessing robustness*

1011 The wake entrophy squares the vorticity, so the per-encounter high-percentile amplitude clip of the ω_z
 1012 preprocessing could in principle shape the primary endpoint. We test this on the frozen encoders, with no
 1013 retraining, by recomputing the readout-frame wake-entrophy target under three training-only clip treatments of
 1014 the raw field and re-fitting the frozen-latent-to-wake ridge probe per family: no amplitude clip, the production
 1015 per-encounter $p_{99.99}$ clip, and a single leak-free $p_{99.99}$ threshold pooled over the training encounters. The
 1016 predictive latent leads the wake on test_b under all three (table 9), with the largest shift in the predictive-minus-
 1017 best-baseline advantage relative to the production clip being 0.07 in R^2 . The leak-free training-global clip gives
 1018 the largest predictive advantage, so the per-encounter clip is not the source of the wake result.

1019 *Return to the baseline orbit within the encounter window*

1020 The load-bearing caveat of §4.3, that the simulation gust trajectories do not fully return to the baseline orbit
 1021 within the 120-frame encounter, is a consequence of the gust-release cadence rather than a strength-dependent
 1022 failure to recover. Figure 14 holds the core diameter fixed at $D = 1.0$ with $G < 0$, averages the distance
 1023 to the baseline orbit over each case’s encounters, and sweeps $|G|$. The departure amplitude near impact is
 1024 nearly independent of gust strength (peak distance 3.8 to 4.2 orbit diameters across $|G| \in [0.5, 3.0]$), every
 1025 trajectory is still contracting toward the orbit at the end of the window, and the strongest gust ends closest rather
 1026 than farthest. The residual offset at the window’s end is therefore set by the $\sim 4t/c$ of post-impact relaxation
 1027 the cadence allows and by the core diameter, a larger core leaving a more persistent wake, not by $|G|$. This
 1028 window-truncated return is common to all three encoders and does not affect the relative drift comparison of
 1029 §4.3, which contrasts the predictive and reconstructive rollouts over the same window.

1030 **Appendix B. Sparse wall-pressure state estimation**

1031 At deployment the vorticity field is not available; only the wall pressure is. We ask how much of the latent state,
 1032 the gust parameters, and the recovered flow field can be read from a sparse set of K wall-pressure taps, and
 1033 through which representation. The input is the span-averaged pressure on the 192 surface taps; the pressure-to-
 1034 latent map is a kernel-ridge regressor with a radial-basis kernel on the flattened K -sensor by window feature
 1035 vector, fitted on the training encounters and applied unchanged on the held-out set. Two windows are reported.
 1036 The primary feasibility comparison uses a window of $W = 30$ frames ending at the readout instant (impact plus
 1037 $H = 16$, the closure-matrix primary endpoint), so the pressure is available up to but not beyond the readout;
 1038 this is a feasibility comparison across representations, not a closed controller. The causal claim of §4.7 instead
 1039 uses a strictly pre-impact window of the same length ending at the impact frame, so the pressure ends before
 1040 the gust reaches the leading edge: under that window the family ordering is preserved and in fact strengthens
 1041 (predictive 0.82 against 0.77 for the reconstructive autoencoder and 0.22 for the linear basis).




Figure 14. Diameter-controlled return to the baseline orbit. Distance of the encoded gust trajectory to the no-gust orbit (predictive latent, in orbit diameters; grey band is the orbit thickness), at fixed core diameter $D = 1.0$ and $G < 0$, averaged over each case’s encounters, for five gust strengths. The departure amplitude is nearly independent of $|G|$ and all strengths are still contracting when the next gust is released at +79 frames, so the non-return is a window-length effect, not a strength-dependent one.

Target-blind placement

The placement we use throughout is qDEIM, the QR-pivot optimum of the wall-pressure modes (Manohar *et al.* 2018). It is target-blind and family-independent: it places taps to reconstruct the pressure field, with no knowledge of the latent it is later asked to recover, and the same eight taps are used for every representation. The cross-family recoverability ordering of §4.7 therefore cannot be an artefact of a bespoke latent-aware sensor selection. A target-aware greedy selection (target-conditioned sensor importance, which adds at each step the tap that most improves recovery of the leading latent component from the already-selected taps) is also available, and its only measurable advantage over qDEIM is in the scarce-sensor regime; by $K = 8$ the wall pressure is information-rich enough that the placement choice no longer separates the families, which is why we report the target-blind optimum.

State recovery and its anisotropy

Over the feasibility window, with $K = 8$ taps the held-out test_b state R^2 orders the families as in the main text (§4.7): predictive 0.78, reconstructive autoencoder 0.66, variational baseline 0.51, linear basis 0.34. The variance-weighted R^2 is reported alongside the mean canonical correlation between the estimated and true latent, because the two expose different things. The linear basis has the highest canonical correlation (0.64 against 0.56 for the predictive) yet the lowest state R^2 : the wall pressure aligns with many POD directions but those are largely the low-variance ones, so a high directional agreement recovers little of the state’s energy. The predictive latent has the lower canonical correlation but recovers the high-variance directions that carry the state, which is the quantity a downstream probe reads.

Recovered field is a different axis from recovered state

Decoding the pressure-estimated latent through each family’s frozen decoder gives a recovered flow field (figure 15). Scored by structural similarity, the predictive latent reaches 0.57, the linear basis 0.58, and the reconstructive 0.48: the predictive latent beats the reconstructive but ties the linear basis, and the tie is not significant (one-sided sign test $p = 0.860$). The tie is a decoder property, not a state-recovery one. The decode ceiling, the similarity of the field decoded from the *true* encoded latent with no pressure step, is highest for the linear basis (0.71 against 0.62 for the predictive and 0.54 for the reconstructive), so the linear basis reaches field parity through its decoder’s reconstruction ceiling while having the weakest state recovery of the three. The pressure-to-state gap, the ceiling minus the pressure-recovered similarity, is correspondingly largest for the linear basis (0.13 against 0.05 for the predictive): the part of the field error attributable to imperfect pressure-to-state recovery is smallest for the predictive latent. The volume relative error and peak-vorticity retention agree with the similarity ordering at the impingement scale (predictive 0.92 relative error and 0.92 peak retention, against 0.78 for the linear basis). The state-recovery winner is therefore not the field-recovery winner (0 of the two orderings coincide), and the clean comparison is the state recovery, where the predictive latent is unambiguously the most recoverable.

flow-field recovery, $K=8$ qDEIM taps, impact+16; representative case $G+1.50_D1.50_Y+0.10$ enc 0




Figure 15. Flow recovered from $K = 8$ target-blind (qDEIM) wall-pressure taps on a representative held-out encounter (the typical-difficulty case $G = -1.50$, $D = 1.50$, $Y = +0.10$), one row per family. Columns: the field decoded from the pressure-estimated latent, the decode ceiling (the same decoder applied to the true encoded latent, with no pressure step), and the simulation. The per-panel similarity makes the appendix point visible: the linear basis reaches field parity through the highest decode ceiling, not through better pressure-to-state recovery, where the predictive latent has the smallest ceiling gap. The decoder is blurry by design (§4.6); the comparison is between the columns of the same row.

Parameters and the out-of-distribution boundary

The gust parameters recover directly from the same eight taps, with the wall-normal offset the weakest (Y at $R^2 = 0.51$ on test_b), consistent with Y 's weak resolution throughout the paper and confirming that the limit is the mid-plane vorticity observable and the offset sampling, not the pressure signal. The impact lift read through the latent ($R^2 = 0.43$) is no better than reading it straight off the pressure, in keeping with the main-text point that the wall reads the scalar lift well but the wake structure only through the representation. On the out-of-distribution test_c set ($|G| = 4$) every estimator and every family gives negative state R^2 , the same three-dimensional observability boundary of §5.2 that bounds the rest of the model.

Robustness and a non-result

The frame-0 gust-release clip of the re-release encounters leaves the result unchanged: removing the clipped frame from the pressure window shifts the predictive state recovery by 0.001 in R^2 , within noise. The wall pressure does not recover the wake: a pressure-to-latent-to-frozen-wake-probe route reaches only $R^2 = 0.18$ for the predictive latent, a direct pressure-to-wake regression 0.17, and the predictive advantage over the best reconstructive family is not significant (one-sided sign test $p = 0.140$). The wake is linearly readable from the encoded field (§4.1) but not from this many wall taps over this window; the deployment comparison is the state and the field, not the wake.




Figure 16. The predictive rollout tracks the encounter, shown qualitatively. Predicted observables through representative held-out encounters (rows), for the lift coefficient, the wake entrophy, and the negative circulation (columns). The simulation is the bold reference; the rollout of the predictor, started from the pre-impact context, is the coloured line, and the dashed line is the reconstruction from the encoded latent. The rollout tracks the wake observables through the leading-edge-vortex peak and into recovery before drifting gently at long times. We report this as qualitative tracking, not as a closure R^2 . Numbered glyphs mark the four encounter stages used throughout: baseline, impact, peak load, and recovery.

Appendix C. Supplementary figures

The figures below are referenced from the Results and collected here to keep the main-text figure sequence focused on the narrative steps. They are diagnostic and qualitative: the predictor rollout traces (figure 16), the per-axis breakdown of the paired wake advantage (figure 17), and the decode-ceiling vorticity panels (figure 18).

REFERENCES

- ARRANZ, GONZALO & LOZANO-DURÁN, ADRIÁN 2024 Informative and non-informative decomposition of turbulent flow fields. *Journal of Fluid Mechanics* **1000**, A95.
- ASSRAN, M., DUVAL, Q., MISRA, I., BOJANOWSKI, P., VINCENT, P., RABBAT, M., LECUN, Y. & BALLAS, N. 2023 Self-supervised learning from images with a joint-embedding predictive architecture. In *CVPR*.
- ASSRAN, M. & OTHERS 2025 V-JEPA 2: self-supervised video models. *arXiv preprint arXiv:2506.09985*.
- BALESTRIERO, R. & LECUN, Y. 2025 LeJEPA: provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*.
- BARDES, A., PONCE, J. & LECUN, Y. 2022 VICReg: variance-invariance-covariance regularization for self-supervised learning. In *ICLR*.
- BERKOOZ, GAL, HOLMES, PHILIP & LUMLEY, JOHN L. 1993 The proper orthogonal decomposition in the analysis of turbulent flows. *Annual Review of Fluid Mechanics* **25**, 539–575.




Figure 17. Per-encounter paired improvement $\Delta e = e_{AE} - e_{JEPA}$ in held-out wake-entropy error against the gust strength G , core diameter D , wall-normal offset Y , and the baseline shedding phase at impact ϕ (test_b, $n = 42$); positive Δe means the predictive latent carries the smaller error. Each panel shows the per-encounter values, a locally weighted (LOWESS) trend, and a bootstrap band over encounters. The predictive advantage grows with D (significant) but shows no significant G or Y trend; the shedding-phase axis falls into two clusters because impact is fixed near a single convective instant.




Figure 18. Decoded vorticity at the impact frame on held-out encounters (rows: an interpolation test_b encounter and a $|G| = 4$ test_c encounter) for the simulation and the predictive, reconstructive, and POD decodes (columns). The predictive decoder is blurry by design but localises the impingement; the reconstructive autoencoder collapses toward a near-uniform field on the held-out interpolation condition, while the linear basis is amplitude-accurate. All families degrade in the $|G| = 4$ regime, reported as the observability boundary of §2.1 rather than as a generalisation claim. Field-space similarity is read only as a labelled decode ceiling (§4.6).

- BRUNTON, STEVEN L., NOACK, BERND R. & KOUMOUTSAKOS, PETROS 2020 Machine learning for fluid mechanics. *Annual Review of Fluid Mechanics* **52**, 477–508.
- CONSTANTE-AMORES, C. RICARDO & GRAHAM, MICHAEL D. 2024 Data-driven state-space and Koopman operator models of coherent state dynamics on invariant manifolds. *Journal of Fluid Mechanics* **984**, R9.
- CREMADES, ANDRÉS, MARTÍNEZ-SÁNCHEZ, ÁLVARO, LOZANO-DURÁN, ADRIÁN & VINUESA, RICARDO 2026 X-CAL: Explaining latent causality in physical space for fluid mechanics. *arXiv preprint* ArXiv:2601.03311.
- ELDREDGE, JEFF D. & JONES, ANYA R. 2019 Leading-edge vortices: Mechanics and modeling. *Annual Review of Fluid Mechanics* **51**, 75–104.
- FUKAMI, KAI & ARAKI, RYO 2026 Information-theoretic machine learning for time-varying mode decomposition of separated aerodynamic flows. *AIAA Journal* **64** (2), 605–613.
- FUKAMI, KAI, NAKAO, HIROYA & TAIRA, KUNIHIKO 2024 Data-driven transient lift attenuation for extreme vortex gust–airfoil interactions. *J. Fluid Mech.* **992**, A17.
- FUKAMI, KAI, SMITH, LUKE & TAIRA, KUNIHIKO 2025 Extreme vortex-gust airfoil interactions at Reynolds number 5000. *Phys. Rev. Fluids* **10**, 084703.
- FUKAMI, KAI & TAIRA, KUNIHIKO 2023 Grasping extreme aerodynamics on a low-dimensional manifold. *Nat. Commun.* **14**, 6480.

- 1125 FUKAMI, KAI & TAIRA, KUNIHIKO 2025 Observable-augmented manifold learning for multi-source turbulent
1126 flow data. *J. Fluid Mech.* **1010**, R4.
- 1127 GASPARINO, L., SPIGA, F. & LEHMKUHL, O. 2024 SOD2D: A GPU-enabled spectral finite elements method for
1128 compressible scale-resolving simulations. *Computer Physics Communications* **297**, 109067.
- 1129 HA, DAVID & SCHMIDHUBER, JÜRGEN 2018 Recurrent world models facilitate policy evolution. In *Advances in*
1130 *Neural Information Processing Systems 31 (NeurIPS)*. ArXiv:1803.10122.
- 1131 JONES, A. R., CETINER, O. & SMITH, M. J. 2022 Physics and modeling of large flow disturbances: discrete gust
1132 encounters for modern air vehicles. *Annu. Rev. Fluid Mech.* **54**, 469–493.
- 1133 KELSHAW, DANIEL & MAGRI, LUCA 2024 Proper latent decomposition. *arXiv preprint arXiv:2412.00785* .
- 1134 KOSHIKAWA, RYO, ARAKI, RYO, LIU, QIONG & FUKAMI, KAI 2026 Convolutional causal learning for aerodynamic
1135 flows. *arXiv preprint ArXiv:2601.19104*.
- 1136 LECUN, YANN 2022 A path towards autonomous machine intelligence. *Tech. Rep.*. OpenReview, version 0.9.2.
- 1137 LEE, KOOKJIN & CARLBERG, KEVIN T. 2020 Model reduction of dynamical systems on nonlinear manifolds
1138 using deep convolutional autoencoders. *Journal of Computational Physics* **404**, 108973.
- 1139 LINOT, ALEC J. & GRAHAM, MICHAEL D. 2023 Dynamics of a data-driven low-dimensional model of turbulent
1140 minimal Couette flow. *Journal of Fluid Mechanics* **973**, A42.
- 1141 LIU, ZHECHENG, BECKERS, DIEDERIK & ELDREDGE, JEFF D. 2025 Model-based reinforcement learning for
1142 control of strongly-disturbed unsteady aerodynamic flows. *AIAA Journal* ArXiv:2408.14685.
- 1143 LUSCH, BETHANY, KUTZ, J. NATHAN & BRUNTON, STEVEN L. 2018 Deep learning for universal linear embeddings
1144 of nonlinear dynamics. *Nature Communications* **9**, 4950.
- 1145 MAES, P.-A., GARRIDO, Q., ASSRAN, M., BARDES, A., RABBAT, M., LECUN, Y. & BALLAS, N. 2026 LeWM:
1146 latent world models for self-supervised video prediction. *arXiv preprint arXiv:2603.19312* .
- 1147 MANOHAR, KRITHIKA, BRUNTON, BINGNI W., KUTZ, J. NATHAN & BRUNTON, STEVEN L. 2018 Data-driven
1148 sparse sensor placement for reconstruction: demonstrating the benefits of exploiting known patterns. *IEEE*
1149 *Control Systems Magazine* **38** (3), 63–86.
- 1150 MARTÍNEZ-SÁNCHEZ, ÁLVARO, LÓPEZ, ESTEBAN, LE CLAINCHE, SOLEDAD, LOZANO-DURÁN, ADRIÁN,
1151 SRIVASTAVA, ANKIT & VINUESA, RICARDO 2023 Causality analysis of large-scale structures in the flow
1152 around a wall-mounted square cylinder. *Journal of Fluid Mechanics* **967**, A1.
- 1153 MOHAMED, ABDULGHANI, MARINO, MATTHEW, WATKINS, SIMON, JAWORSKI, JUSTIN & JONES, ANYA 2023 Gusts
1154 encountered by flying vehicles in proximity to buildings. *Drones* **7** (1), 22.
- 1155 MOTOORI, Y. & GOTO, S. 2019 Generation mechanism of a hierarchy of vortices in a turbulent boundary layer.
1156 *J. Fluid Mech.* **865**, 1085–1109.
- 1157 MURATA, TAKAAKI, FUKAMI, KAI & FUKAGATA, KOJI 2020 Nonlinear mode decomposition with convolutional
1158 neural networks for fluid dynamics. *Journal of Fluid Mechanics* **882**, A13.
- 1159 ODAKA, HIROTO, LOPEZ-DORIGA, BARBARA & TAIRA, KUNIHIKO 2026 Vortical similarities across laminar and
1160 turbulent extreme gust encounters. *J. Fluid Mech.* **1031**, R3.
- 1161 SCHMID, PETER J. 2010 Dynamic mode decomposition of numerical and experimental data. *Journal of Fluid*
1162 *Mechanics* **656**, 5–28.
- 1163 SMITH, LUKE, FUKAMI, KAI, SEDKY, GIRGUIIS, JONES, ANYA & TAIRA, KUNIHIKO 2024 A cyclic perspective on
1164 transient gust encounters through the lens of persistent homology. *J. Fluid Mech.* **980**, A18.
- 1165 SOBAL, V., ZHANG, W., CHO, K., BALESTRIERO, R., RUDNER, T. G. J. & LECUN, Y. 2025 Learning from
1166 reward-free offline data: a case for planning with latent dynamics models. *arXiv preprint arXiv:2502.14819*
1167 .
- 1168 SOLERA-RICO, ALBERTO, GARCÍA-CASPUEÑAS, PATRICIA, SANMIGUEL VILA, CARLOS & DISCETTI, STEFANO 2025
1169 The balance between compactness and forecast accuracy of data-driven latent-space reduced-order models
1170 in controlled wake flows. Under review, *Journal of Fluid Mechanics*.
- 1171 SOLERA-RICO, ALBERTO, SANMIGUEL VILA, CARLOS, GÓMEZ-LÓPEZ, M., WANG, Y., ALMASHJARY, A., DAWSON,
1172 S. T. M. & VINUESA, R. 2024 β -variational autoencoders and transformers for reduced-order modelling of
1173 fluid flows. *Nat. Commun.* **15**, 1361.
- 1174 TAYLOR, G. I. 1918 On the dissipation of eddies. *Reports and Memoranda, Advisory Committee for Aeronautics*
1175 (598).
- 1176 TRAN, JONATHAN, YEH, CHI-AN & TAIRA, KUNIHIKO 2026 Using optimal transport aligned latent embeddings
1177 for separated flow analysis. *J. Fluid Mech.* **1027**, A24.
- 1178 ZAMANI ASHTIANI, S. & FUKAMI, KAI 2025 Data-driven time-dependent bases for turbulent airfoil wake–
1179 extreme vortex gust interactions. *arXiv preprint ArXiv:2512.09523*.
- 1180 ZHONG, Y., AMIRI-MARGAVI, A., BABAEE, H. & TAIRA, KUNIHIKO 2025 Optimally time-dependent modes of
1181 vortex gust–airfoil interactions. *Journal of Fluid Mechanics* **1006**, A18.