

Autoregressive transformer predictor P_ϕ (causal, 6 blocks)

Δt embeddings RoPE positions

$\{\mathbf{z}_{t-T+1}, \dots, \mathbf{z}_t\}$ sequence length T

$\mathbf{c} = (G, D, Y)$
gust parameters

Block 1

Multi-head self-attention (16 heads, $d_{\text{model}}=384$)
RoPE on Q, K ; causal mask

FFN (hidden 1536, dropout 0.1)
AdaLN-Zero modulation by $\mathbf{c}$

Block 2

Multi-head self-attention (16 heads)
RoPE on Q, K ; causal mask

FFN (hidden 1536, dropout 0.1)
AdaLN-Zero modulation by $\mathbf{c}$

$\vdots$


Block 6

Multi-head self-attention (16 heads)
RoPE on Q, K ; causal mask

FFN (hidden 1536, dropout 0.1)
AdaLN-Zero modulation by $\mathbf{c}$

$\hat{\mathbf{z}}_{t+1} \in \mathbb{R}^d$

Causal mask



future $\rightarrow$ past blocked

Hyperparameters
6 layers, hidden 384
16 attention heads
dropout 0.1
AdamW, lr 5×10^{-4}